

# Parallel Computing Using MPI, OpenMP, CUDA with examples and debugging, tracing, profiling of parallel programs on the Discovery Cluster.

The session is structured as follows:

- Using MPI on Discovery and brief MPI features (10 minutes)
- OpenMP – Programming with Shared Memory (10 minutes)
- NVIDIA CUDA 6.5 SDK (10 minutes)
- MPI/OpenMP brief examples (15 minutes)
- CUDA/OpenCL brief examples (15 minutes)
- Hybrid MPI, OpenMP example runs on non-GPU nodes (15 minutes)
- Hybrid MPI, OpenMP, CUDA example runs on GPU nodes (15 minutes)
- Debugging, Tracing and Profiling with MPE (10 minutes)
- Debugging, Tracing and Profiling with HPCToolKit (10 minutes)
- Questions (10 minutes)

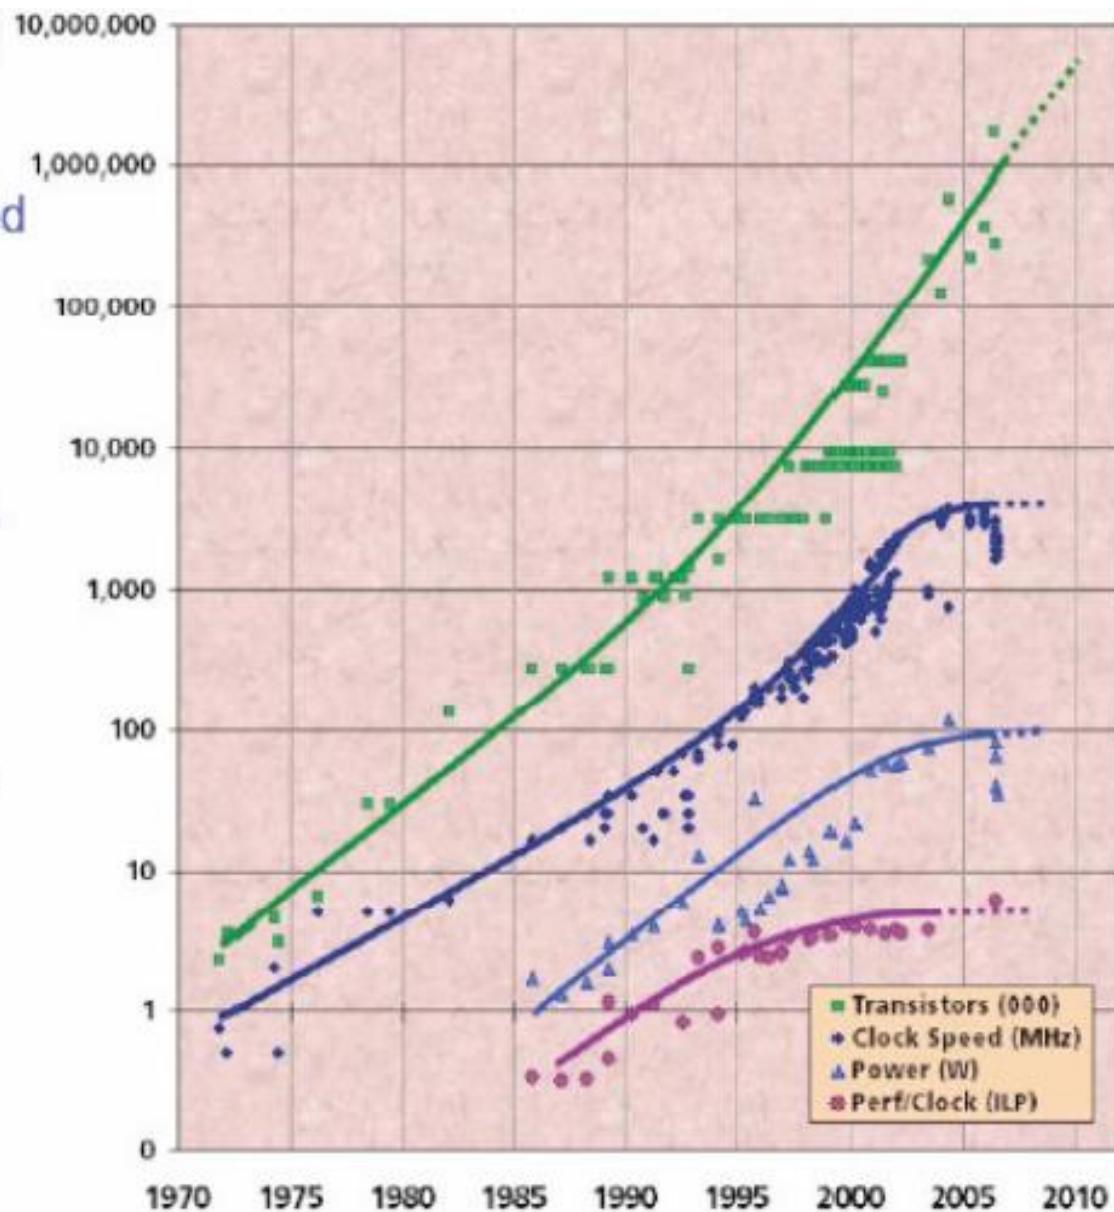
# Why Parallelization ...?

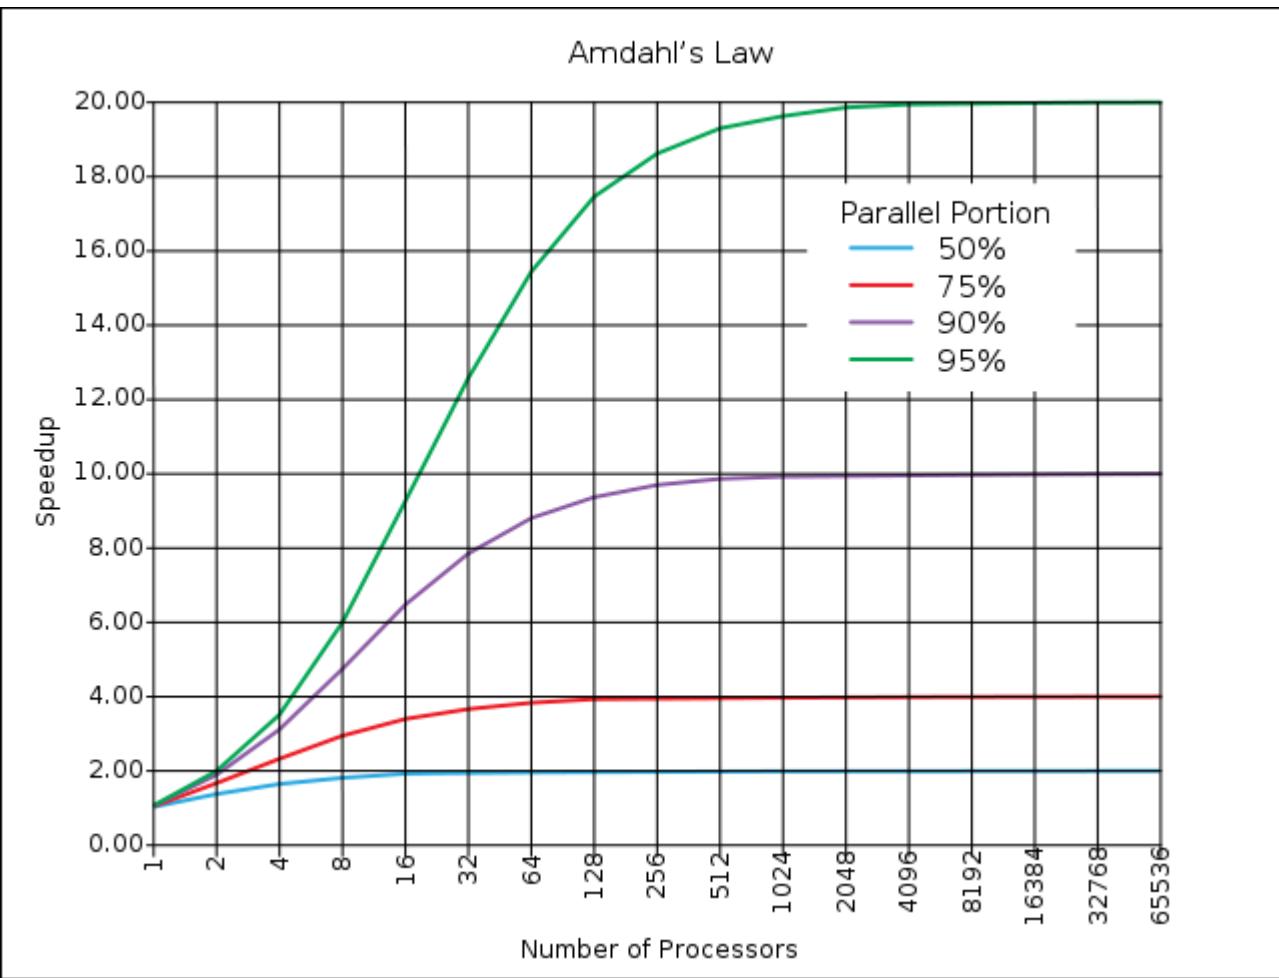
## New Constraints

- 15 years of exponential clock rate growth has ended

## But Moore's Law continues!

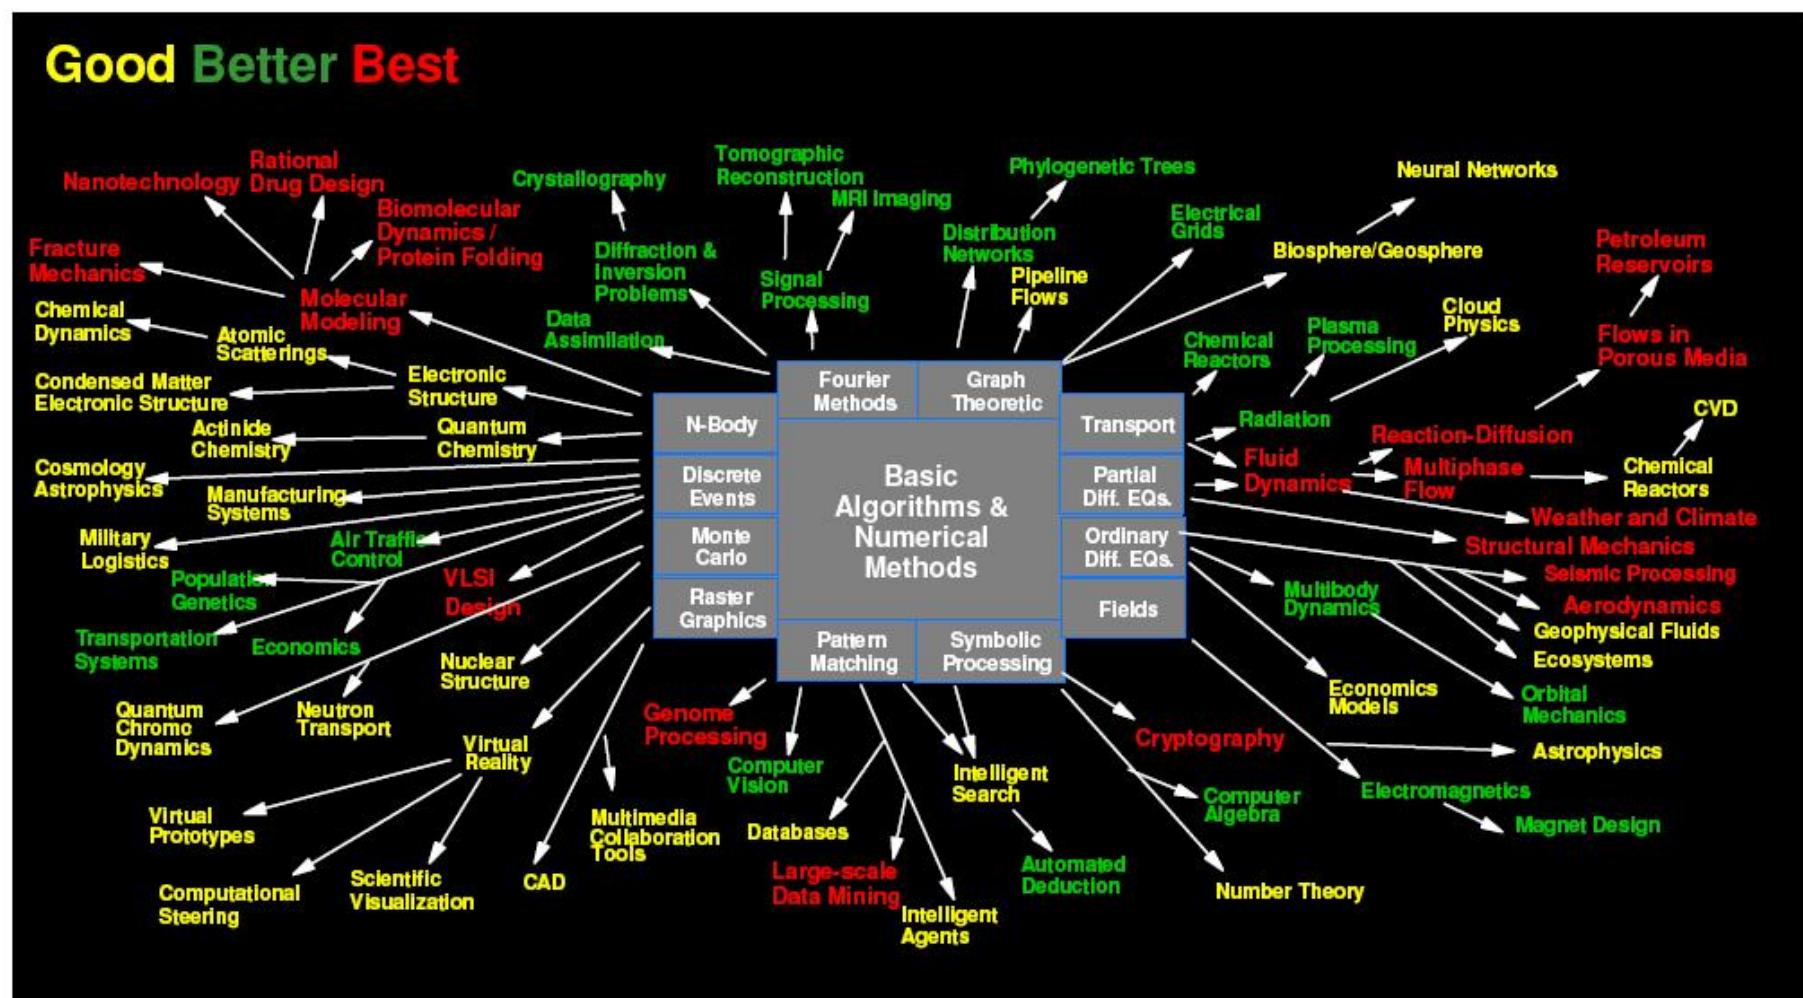
- How do we use all of those transistors to keep performance increasing at historical rates?
- Industry Response: #cores per chip doubles every 18 months *instead of* clock frequency!



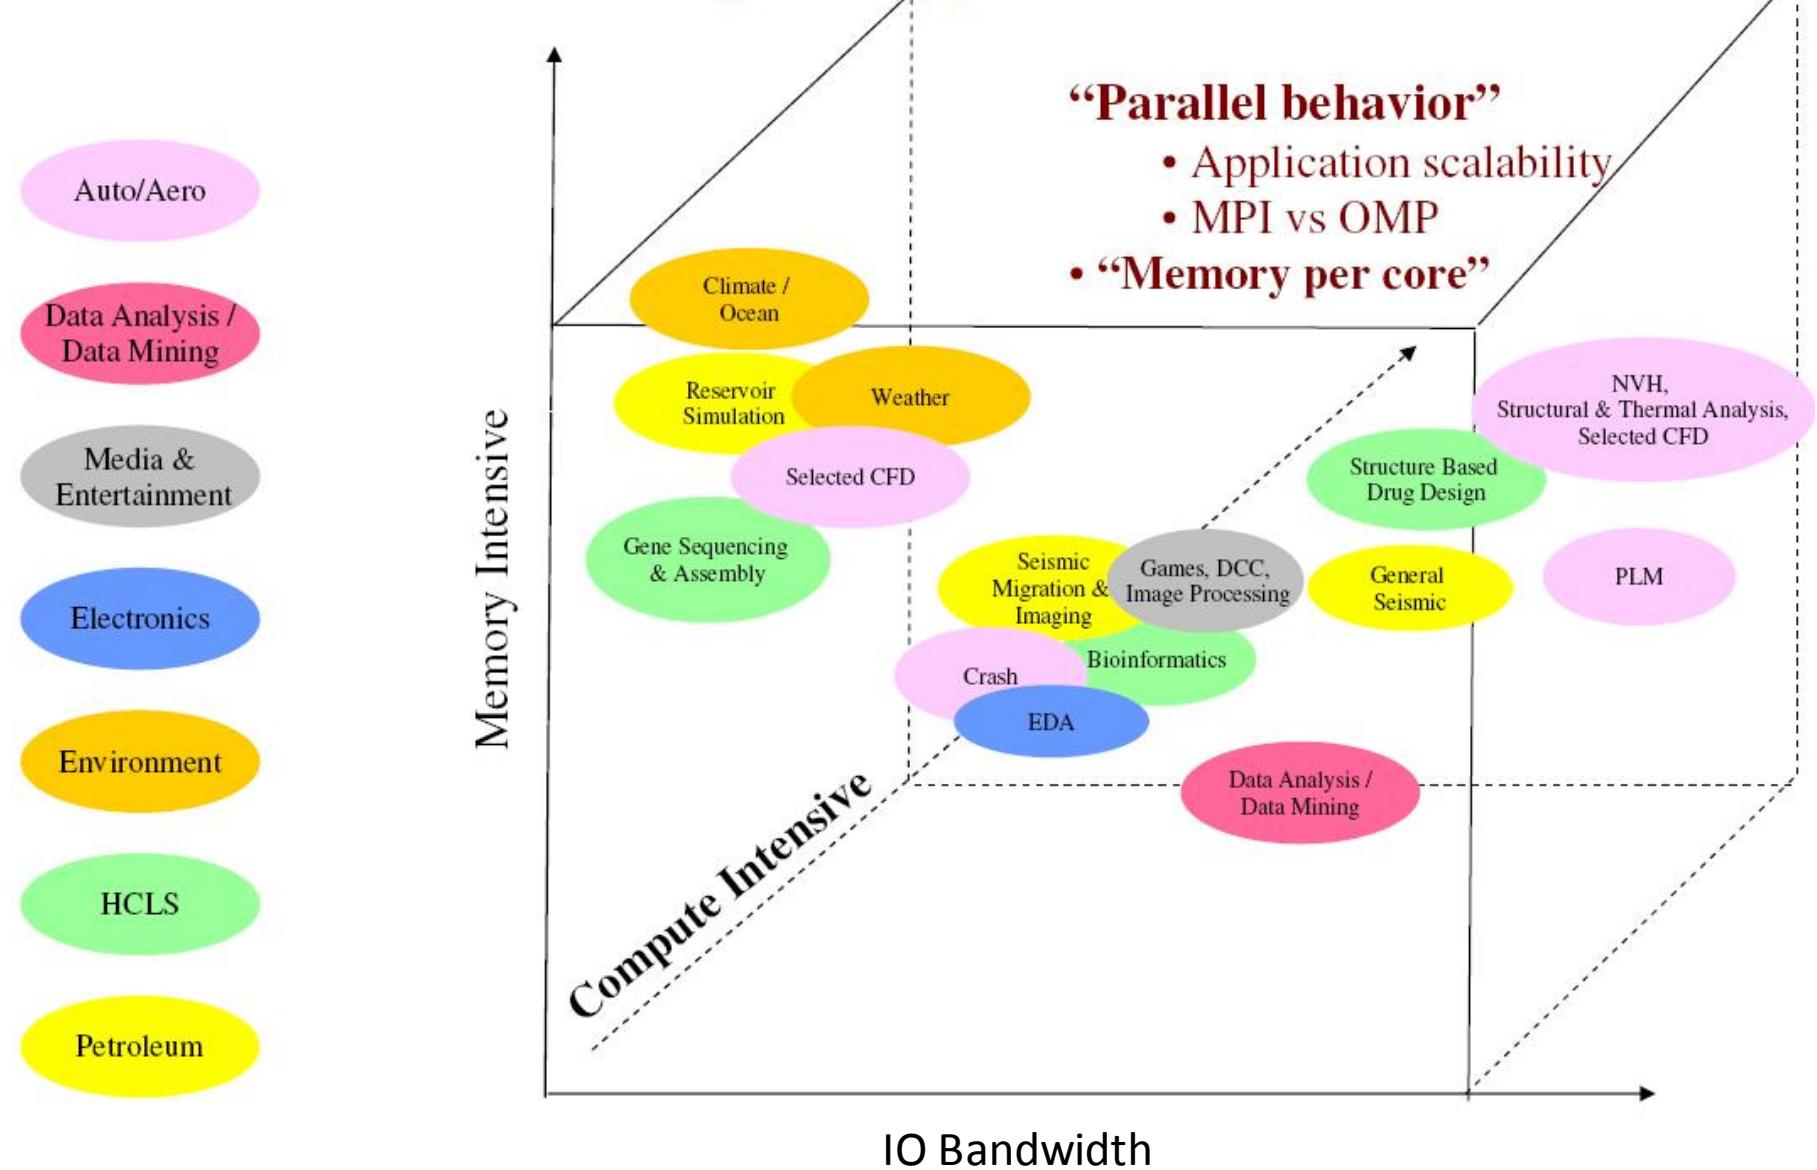


- Amdahl's law states that if  $P$  is the proportion of a program that can be made parallel, and  $(1-P)$  is the proportion that cannot be parallelized, then the maximum speedup that can be achieved by using  $N$  processors =  $P/[(1-P) + (P/N)]$
- 95% of the program can be parallelized, the theoretical maximum speedup using parallel computing would be 20x, no matter how many processors are used

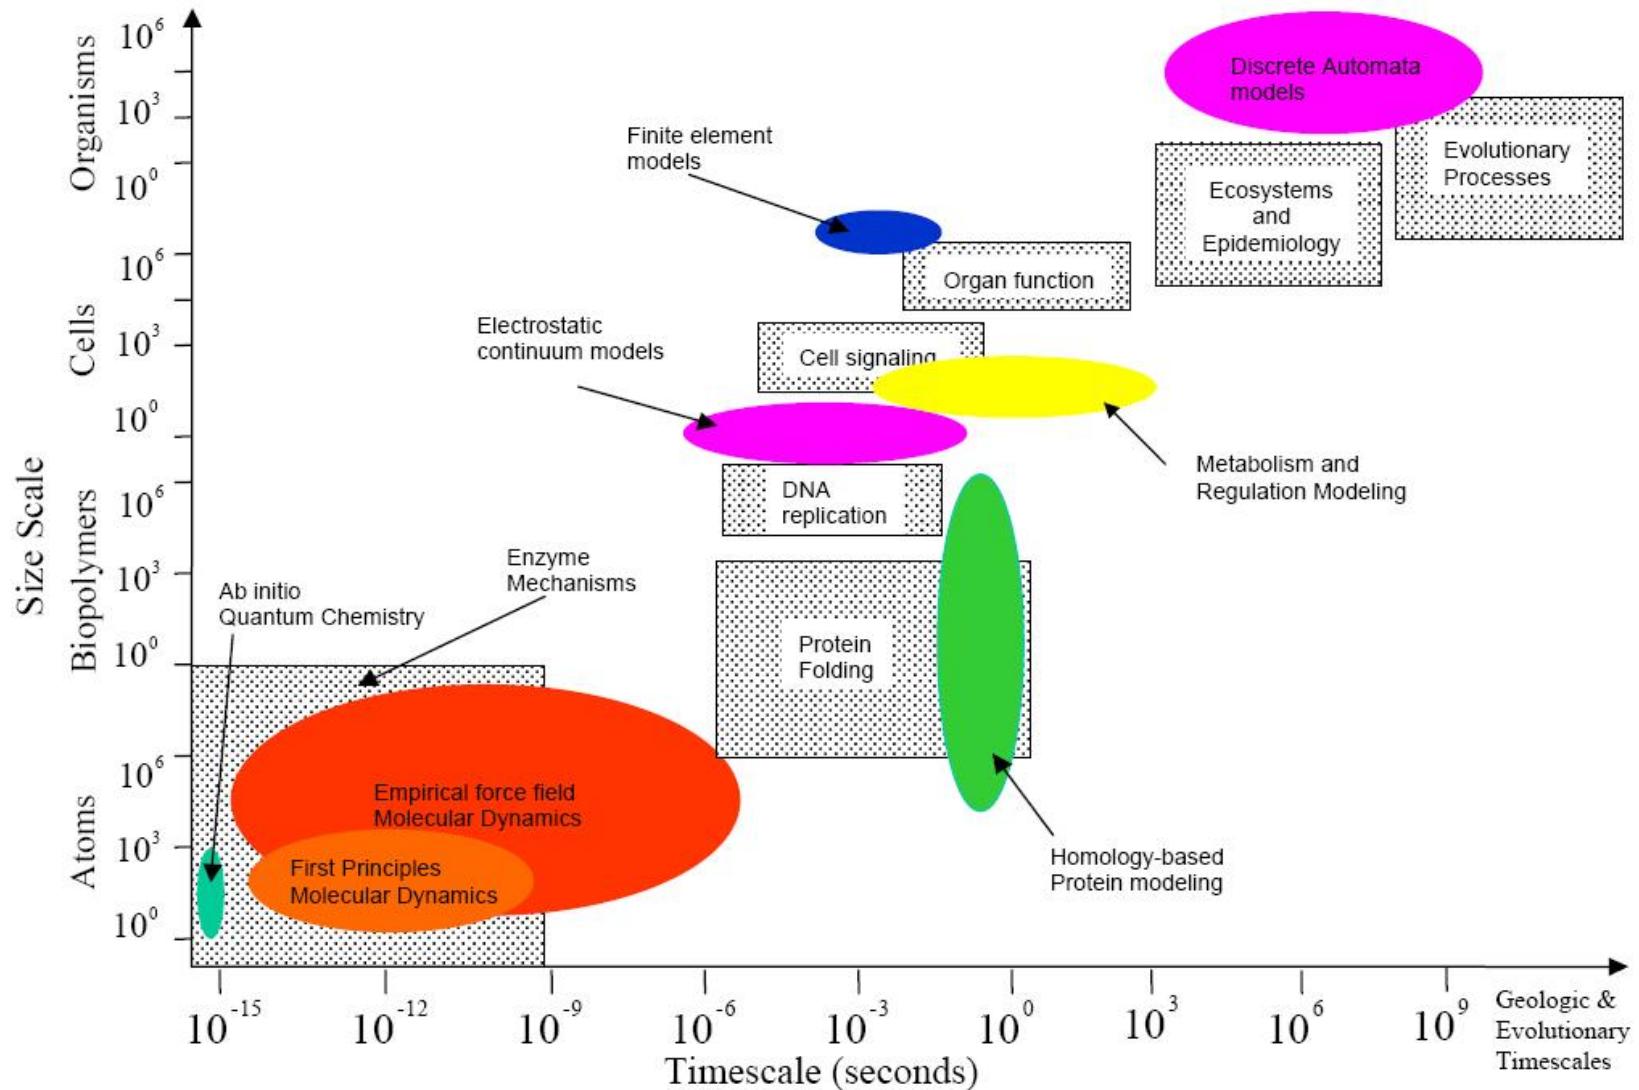
## Many Computational Science Modeling and Simulation Algorithms and Numerical Methods are Massively Parallel



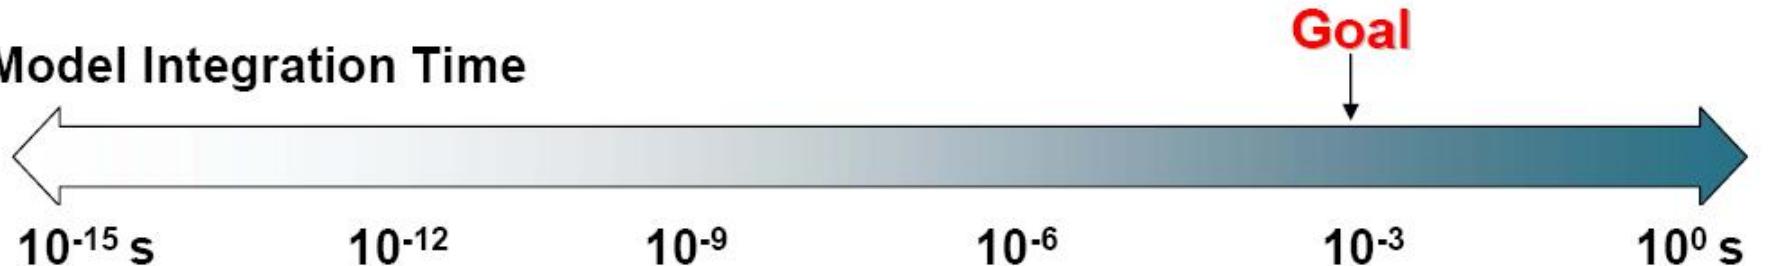
# Platform Positioning – Application Characteristics



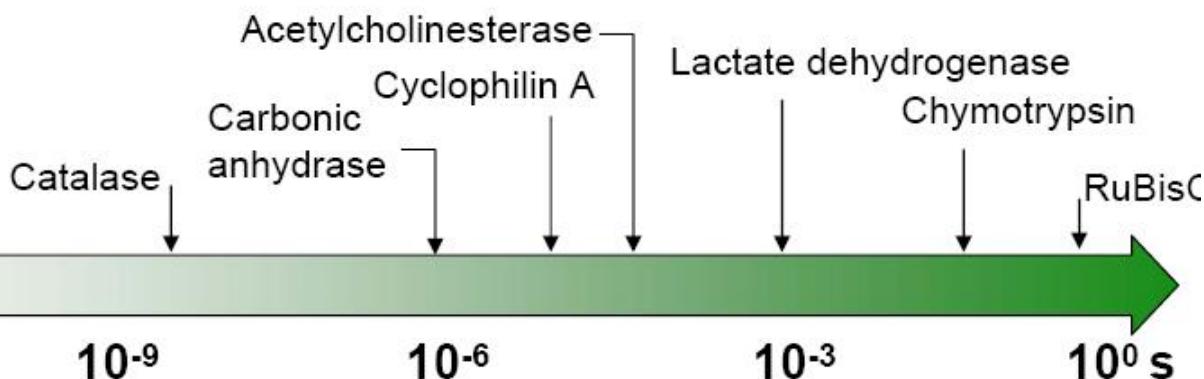
# 24 Orders Magnitude of Spatial and Temporal Range



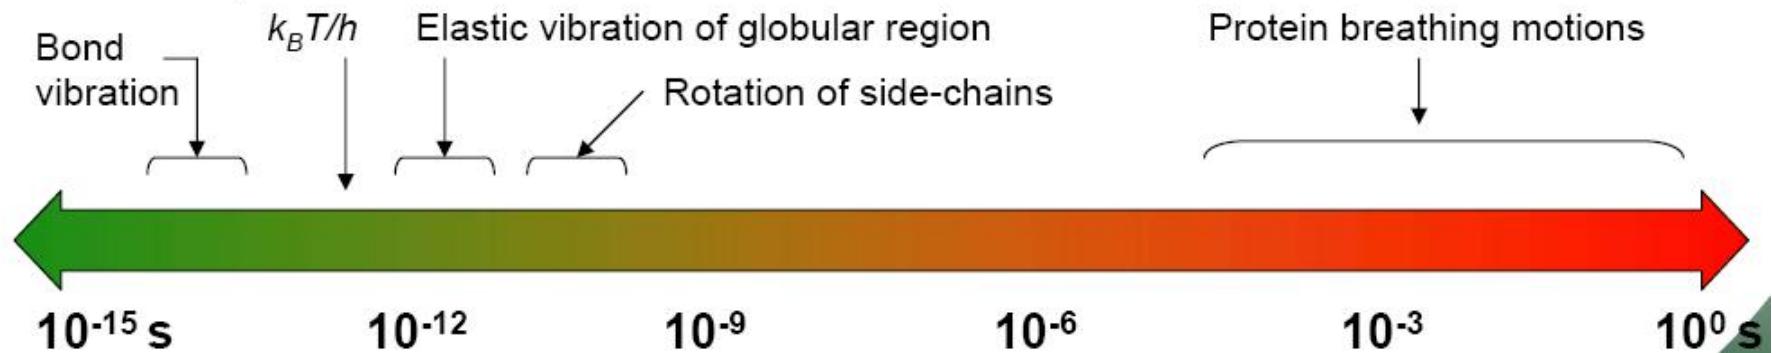
## Model Integration Time



## Enzyme function

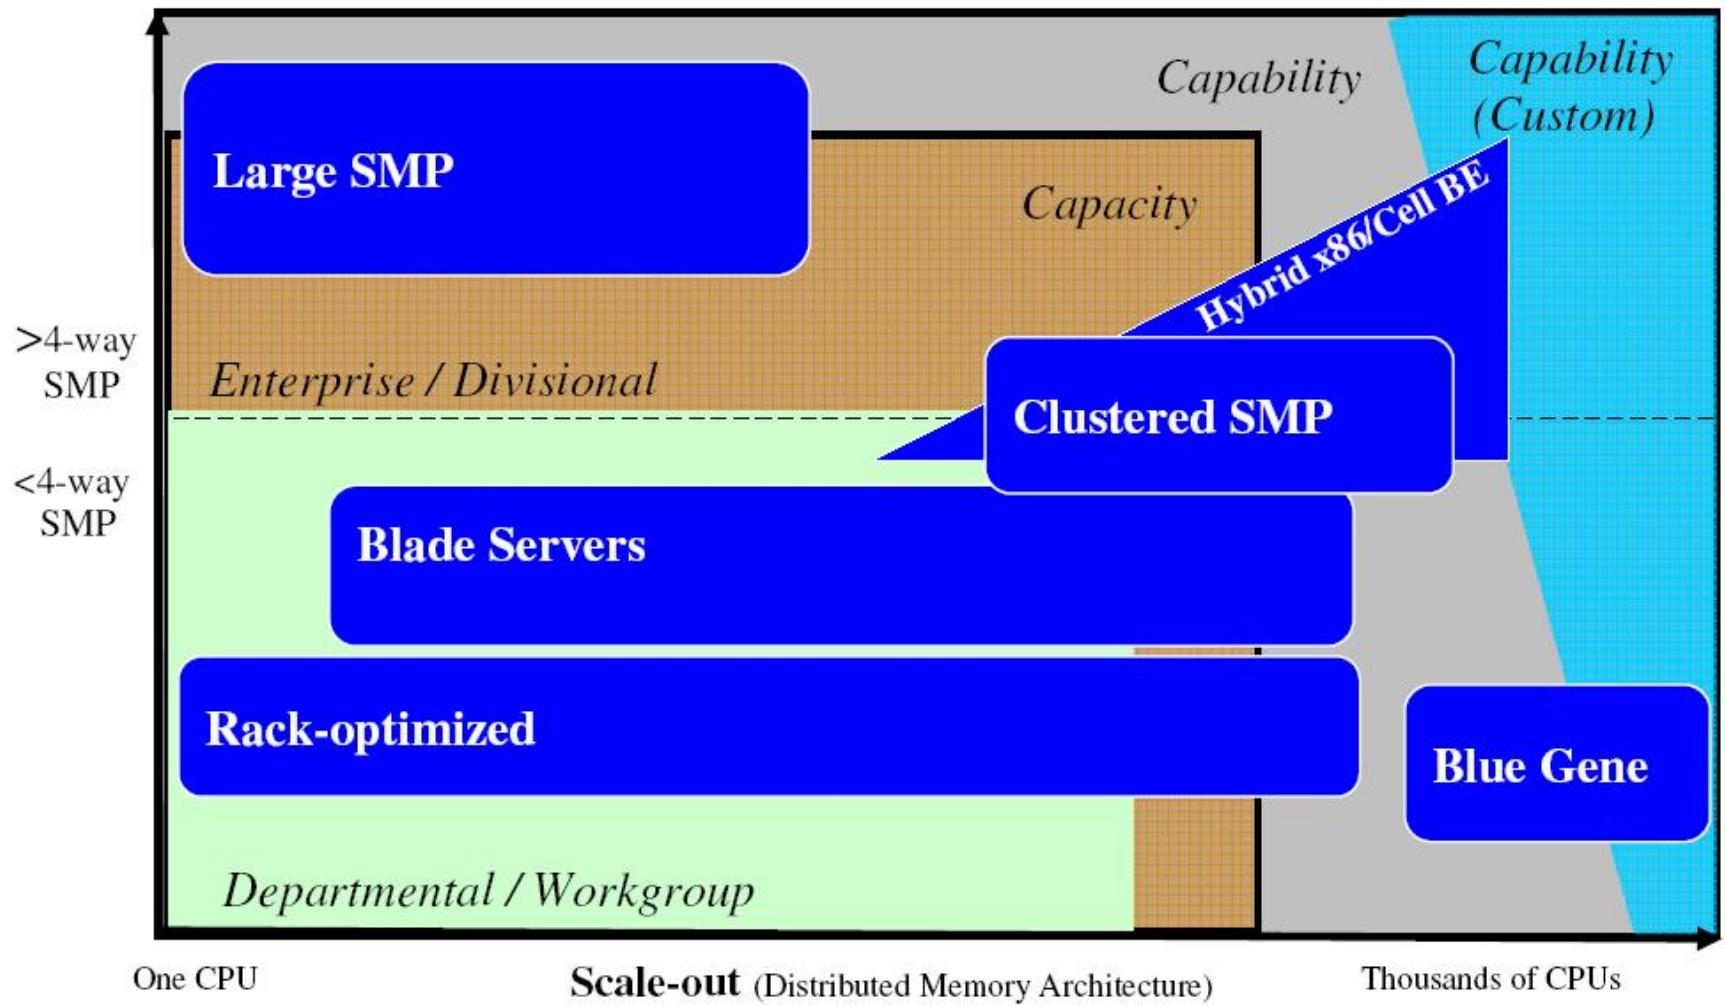


## Protein dynamical events

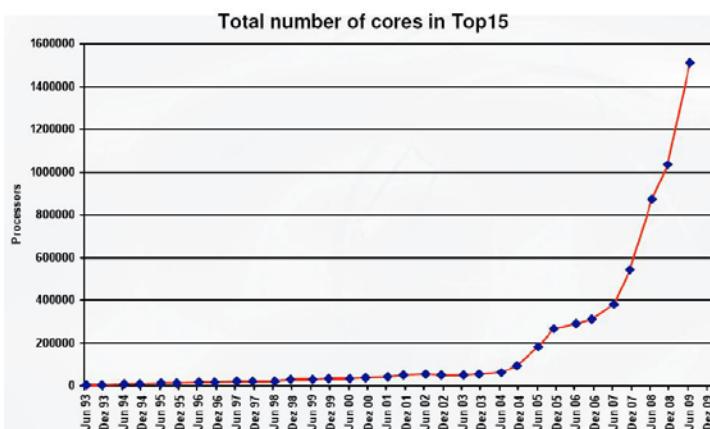
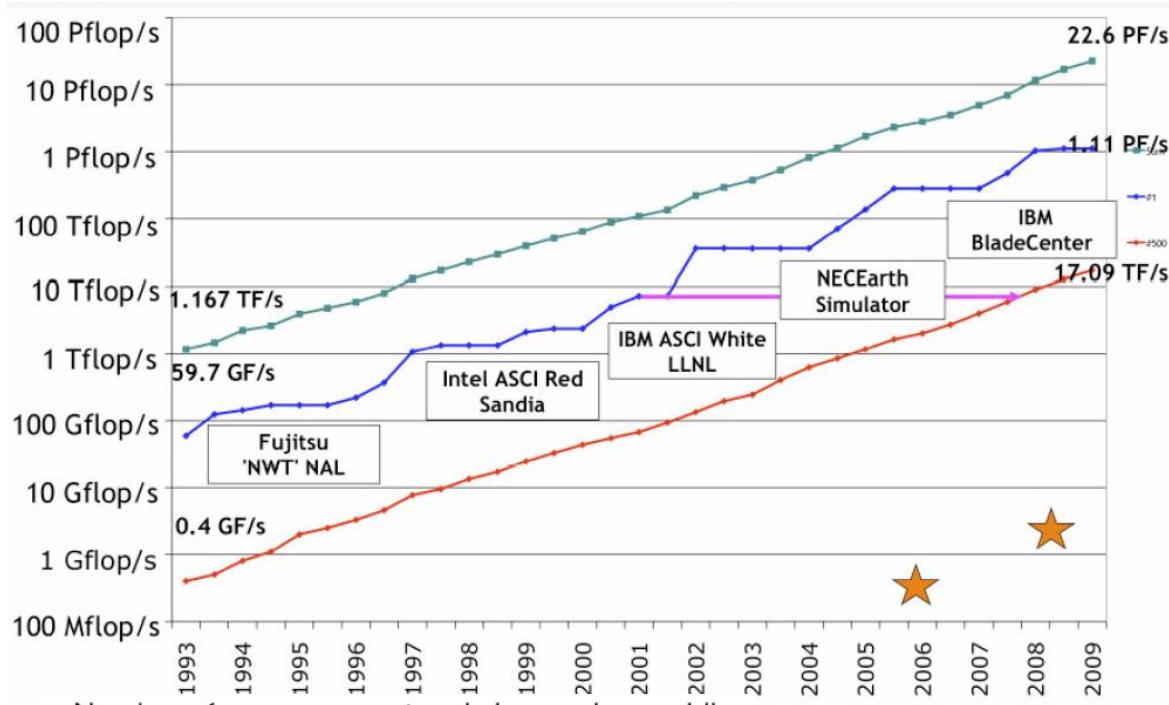


# Platform Positioning - Scalability

**Scale Up**  
(Uniform Memory Architecture)



# HPC Top500 Performance Evolution

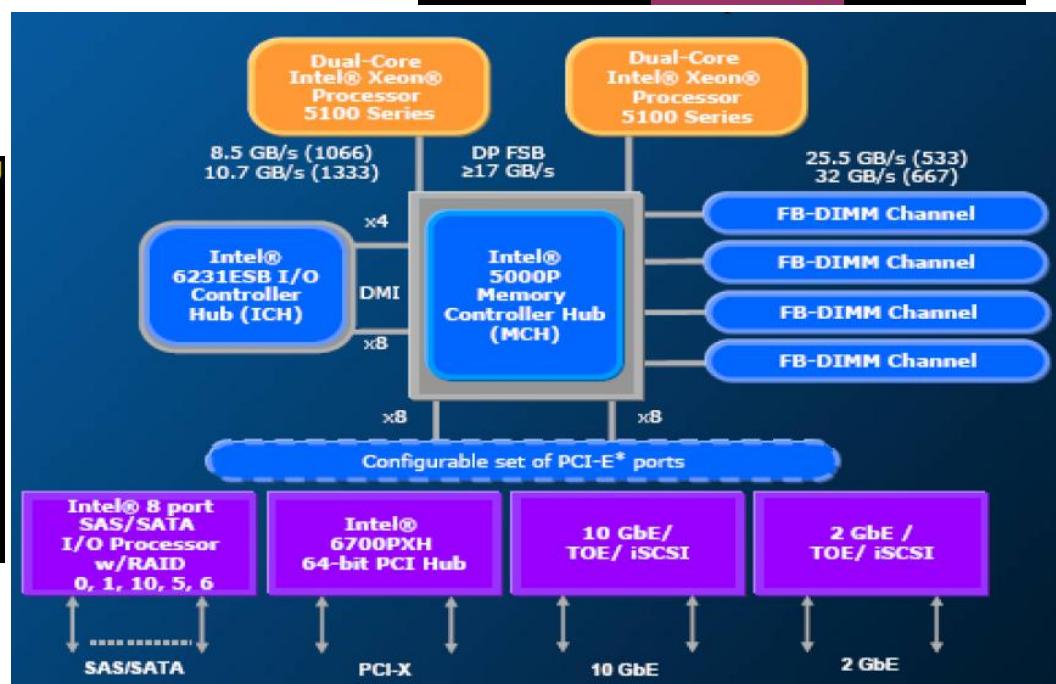
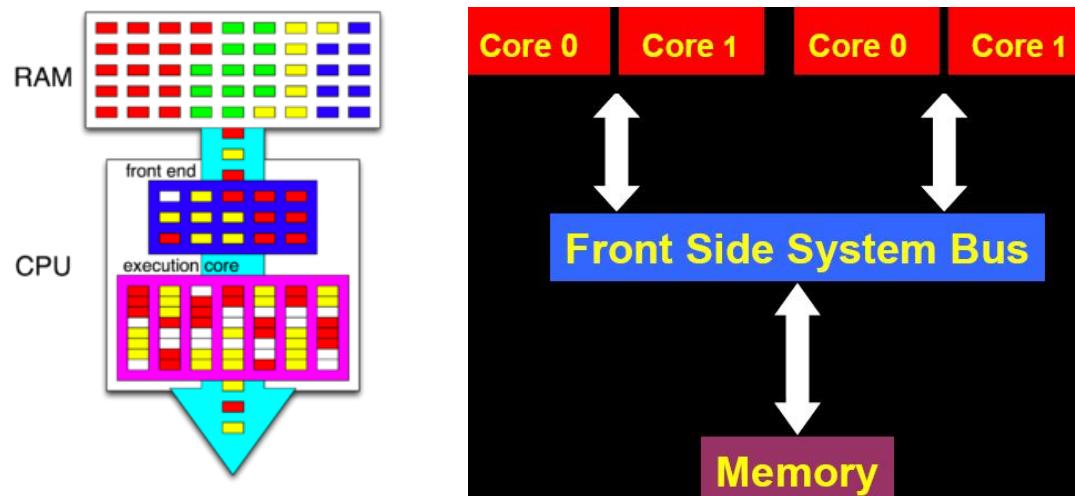
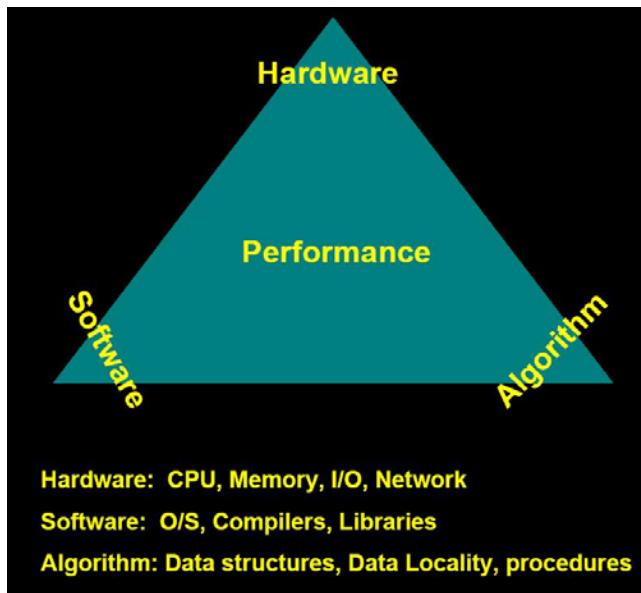


## HPC Application Spectrum

- Bandwidth and processor compute capability assessed
- Applications span the spectrum
- No single industry accepted metric exists



# CPU design, scalability, memory, i/o, overall architectures



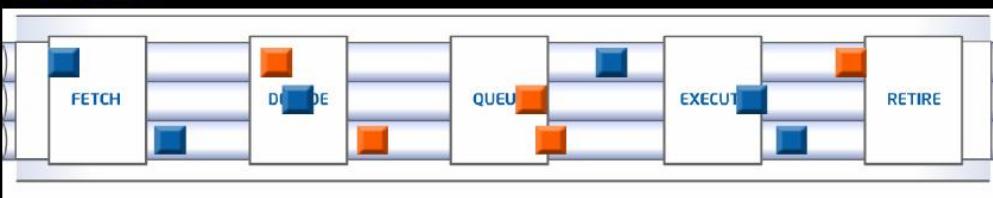
What is important to Software Performance As far as CPU is concerned?

- CPU Speed
- L1/L2 cache size
- L1/L2 Latency
- Execution rate (keeping the processor busy)
- Taking advantage of the Instruction Set
- Support for Threading

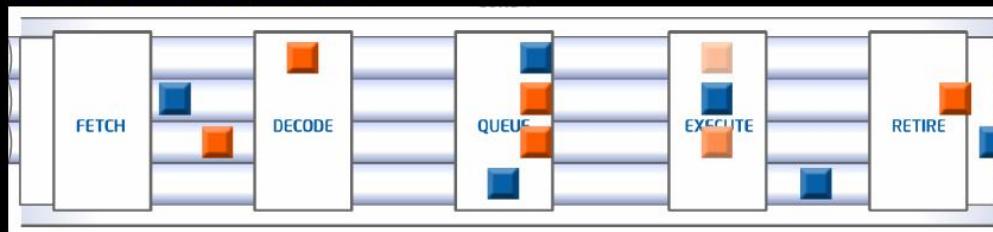
- Instructions per cycle
- Pipeline depth
- SMP and Hyperthreading
- Branch prediction
- Caches with hardware prefetch
- Out of order execution cores
- Instruction level parallelism
- Intra-register vectorization

- **Executes 4 instructions per clock cycle compared to 3 instructions per cycle for NetBurst**

**Net Burst**



**Core Microarchitecture**



## Hyper-Threading

### To improve Single Core performance of

- Multi-threaded Application
- Multi-threaded Operating System
- Single-threaded App in Multi-tasking env

Many levels of parallelism

- Node
- Socket
- Chip
- Core
- Thread
- Register/SIMD
- Multiple instruction pipelines

***MPI***

# Message Passing Interface

# Outline

- Background
- Message Passing
- MPI
  - Group and Context
  - Communication Modes
  - Blocking/Non-blocking
  - Features
  - Programming / issues

# Message Passing

- A process is a program counter and address space.
- Message passing is used for communication among processes.
- Inter-process communication:
  - Type:  
Synchronous / Asynchronous
  - Movement of data from one process's address space to another's

# Synchronous Vs. Asynchronous

- A synchronous communication is not complete until the **message** has been received.
- An asynchronous communication completes as soon as the **message** is on the way.

# What is message passing?

- Data transfer.
- Requires cooperation of sender and receiver
- Cooperation not always apparent in code

# What is MPI?

- A message-passing library specifications:
  - Extended message-passing model
  - Not a language or compiler specification
  - Not a specific implementation or product
- For parallel computers, clusters, and heterogeneous networks.
- Communication modes: *standard, synchronous, buffered, and ready.*
- Designed to permit the development of parallel software libraries.
- Designed to provide access to advanced parallel hardware for
  - End users
  - Library writers
  - Tool developers

# Group and Context (cont.)

- Are two important and indivisible concepts of MPI.
- Group: is the set of processes that communicate with one another.
- Context: it is somehow similar to the frequency in radio communications.
- Communicator: is the central object for communication in MPI. Each communicator is associated with a group and a context.

# Communication Modes

- Based on the type of send:
  - Synchronous: Completes once the acknowledgement is received by the sender.
  - Buffered send: completes immediately, unless if an error occurs.
  - Standard send: completes once the message has been sent, which may or may not imply that the message has arrived at its destination.
  - Ready send: completes immediately, if the receiver is ready for the message it will get it, otherwise the message is dropped silently.

# Blocking vs. Non-Blocking

- Blocking, means the program will not continue until the communication is completed.
- Non-Blocking, means the program will continue, without waiting for the communication to be completed.

# Features of MPI

- General
  - Communications combine context and group for message security.
  - Thread safety can't be assumed for MPI programs.

# Features that are NOT part of MPI

- Process Management
- Remote memory transfer
- Threads
- Virtual shared memory

# Why to use MPI?

- MPI provides a powerful, efficient, and portable way to express parallel programs.
- MPI was explicitly designed to enable libraries which may eliminate the need for many users to learn (much of) MPI.
- Portable !!!!!!!!!!!!!!!
- Good way to learn about subtle issues in parallel computing

Parallelism can be found/exists at different granularities

- Instruction Level

Ex: add instruction executes with multiply instruction

Compiler good at finding this

- Thread Level

Ex: screen redraw function executes with recalculate in spreadsheet

Programmers good at finding this

- Process Level

Ex: Simulation job runs on same machines as spreadsheet

Users good at creating this

Thread Level Parallelism

Programmer generally makes TLP explicit

Compilers can extract threads in regular programs

```
for (i = 0; i < 200; i++)
for(j = 1; j < 20000; j++)
val[i,j] = val[i,j-1] + 1;
```

```
forall(i = 0; i < 200; i++)
for(j = 1; j < 20000; j++)
val[i,j] = val[i,j-1] + 1;
```

Thread Level Parallelism

Synchronization

- Unlike in ILP, flow of data/dependences must be explicit
- ```
while(ptr = ptr->next)
sum += ptr->val;
```

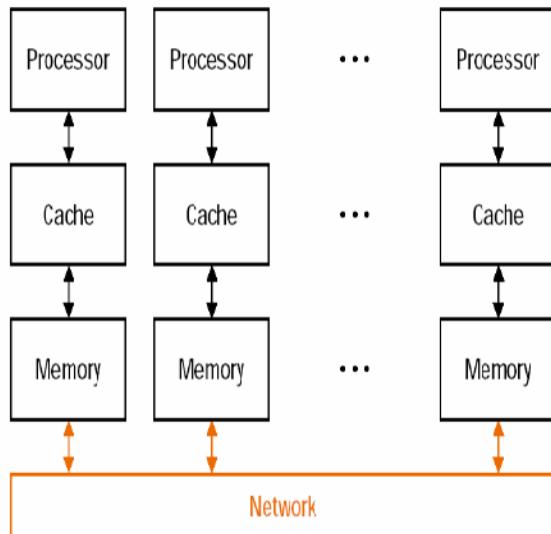
```
while(ptr = ptr->next)
produce(ptr);
produce(NULL);
```

```
while(ptr = consume(ptr))
sum += ptr->val;
```

Communication and Synchronization... (order and flow)

## MPI

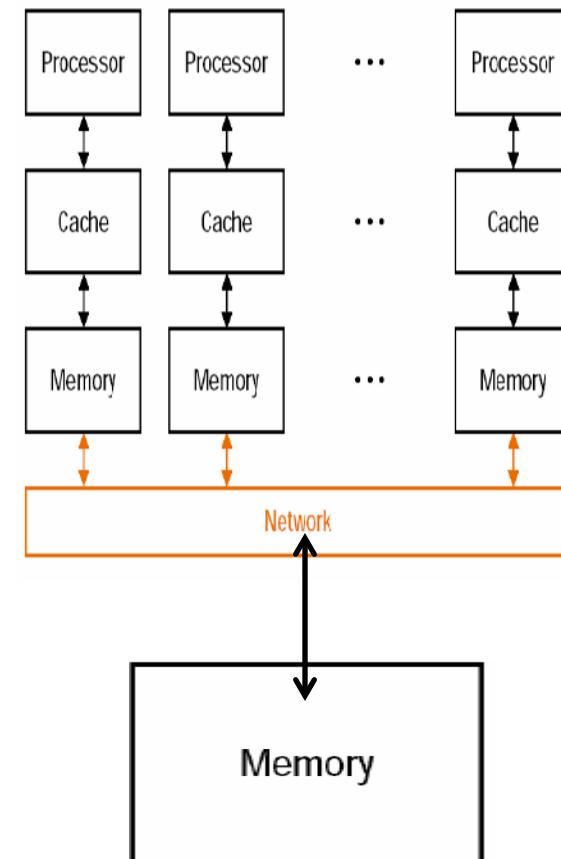
Multiple Processor Organization  
Message Passing/Private Memory



- Threads communicate directly (send, receive)
- Scales relatively well
- No memory coherence problem (for the hardware at least)

## OpenMP

Multiple Processor Organization  
May exist on single chip



# MPI Structure

- All MPI/C/C++ programs must include a header file mpi.h
- All MPI programs must call MPI INT as the first MPI call, to initialize themselves.
- Most MPI programs call MPI COMM SIZE to get the number of processes that are running
- Most MPI programs call MPI COMM RANK to determine their rank, which is a number between 0 and size-1.
- Conditional process and general message passing can take place. For example, using the calls MPI SEND and MPI RECV.
- All MPI programs must call MPI FINALIZE as the last call to an MPI library routine.

# MPI – Point to Consider

- Data types
- Communication - point-to-point and collective
- Timing
- Grouping data for communications
- Communicators and Topologies
- I/O in parallel
- Debugging parallel program
- Performance

# MPI Deadlocks

- Send a large message from process 0 to process 1
  - If there is insufficient storage at the destination, the send must wait for the user to provide the memory space (through a receive)
- What happens with

| Process 0                                    | Process 1                                    |
|----------------------------------------------|----------------------------------------------|
| <code>Send(1)</code><br><code>Recv(1)</code> | <code>Send(0)</code><br><code>Recv(0)</code> |

- This is called “unsafe” because it depends on the availability of system buffers

# MPI – Deadlock Solutions

- Order the operations more carefully:

Process 0

**Send(1)**

**Recv(1)**

Process 1

**Recv(0)**

**Send(0)**

- Use non-blocking operations:

Process 0

**Irecv(1)**

**Irecv(0)**

**Waitall**

Process 1

**Isend(0)**

**Isend(1)**

**Waitall**

# RECAP

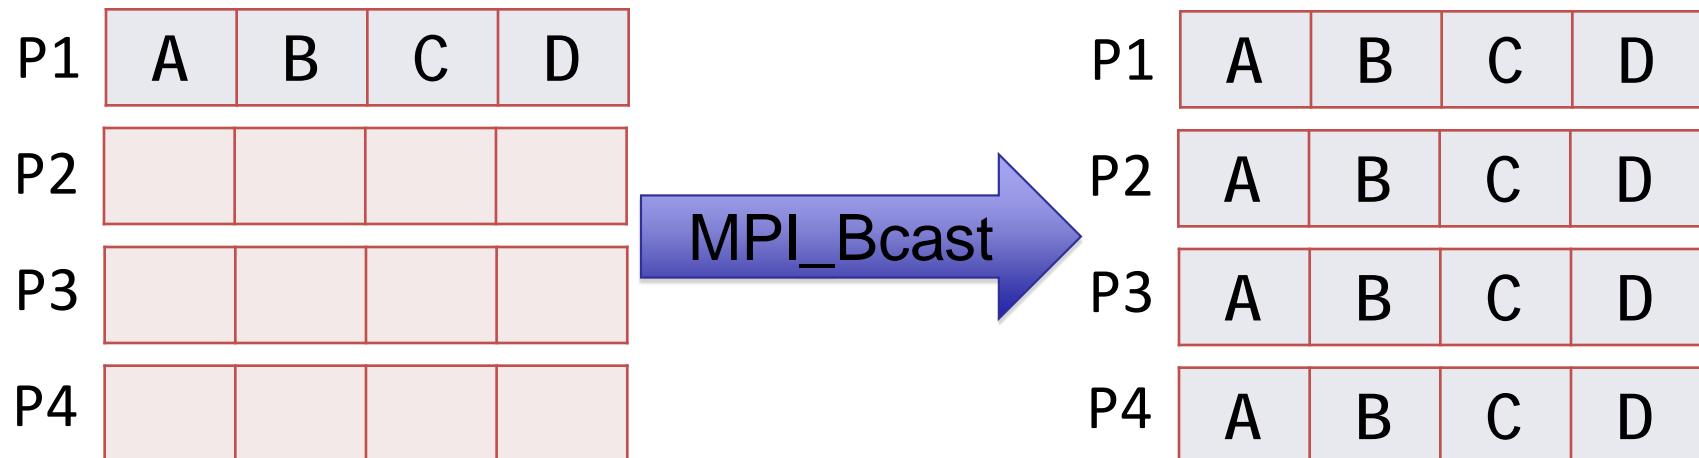
- MPI Communications
- Most important Part of MPI
- RDMA v/s TCP/IP

# Collective communications

- A single call handles the communication between all the processes in a communicator
- There are 3 types of collective communications
  - Data movement (e.g. MPI\_Bcast)
  - Reduction (e.g. MPI\_Reduce)
  - Synchronization (e.g. MPI\_Barrier)

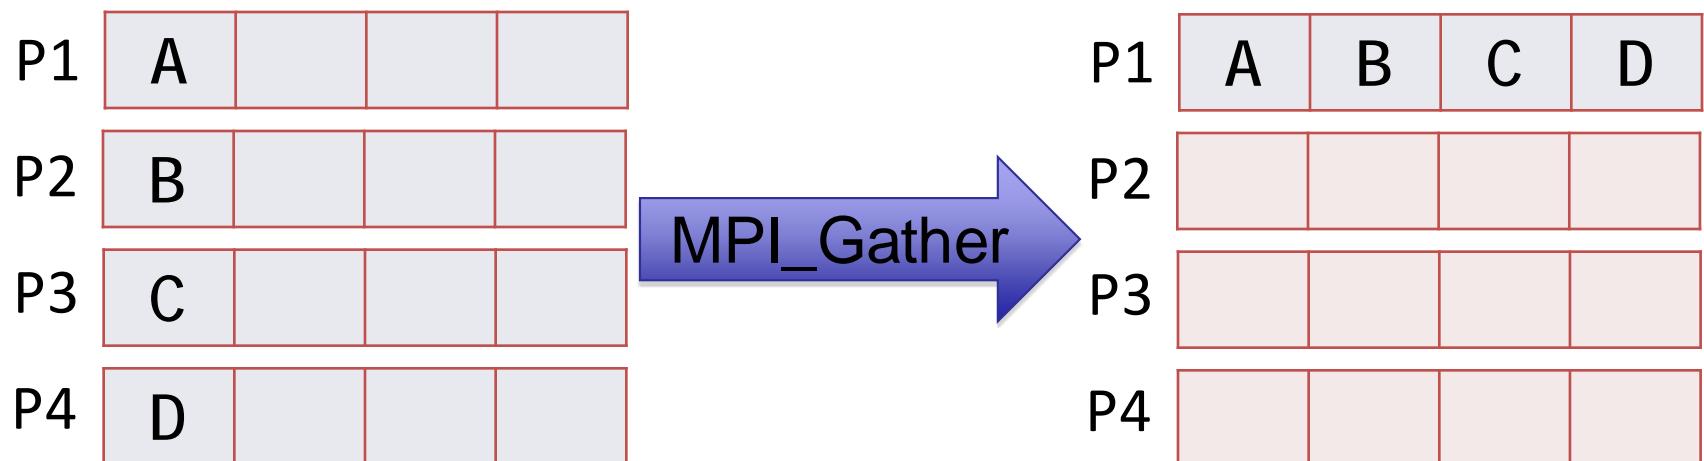
# Broadcast

- int MPI\_Bcast(void \*buffer, int count, MPI\_Datatype datatype, int root, MPI\_Comm comm);
  - One process (root) sends data to all the other processes in the same communicator
  - Must be called by all the processes with the same arguments



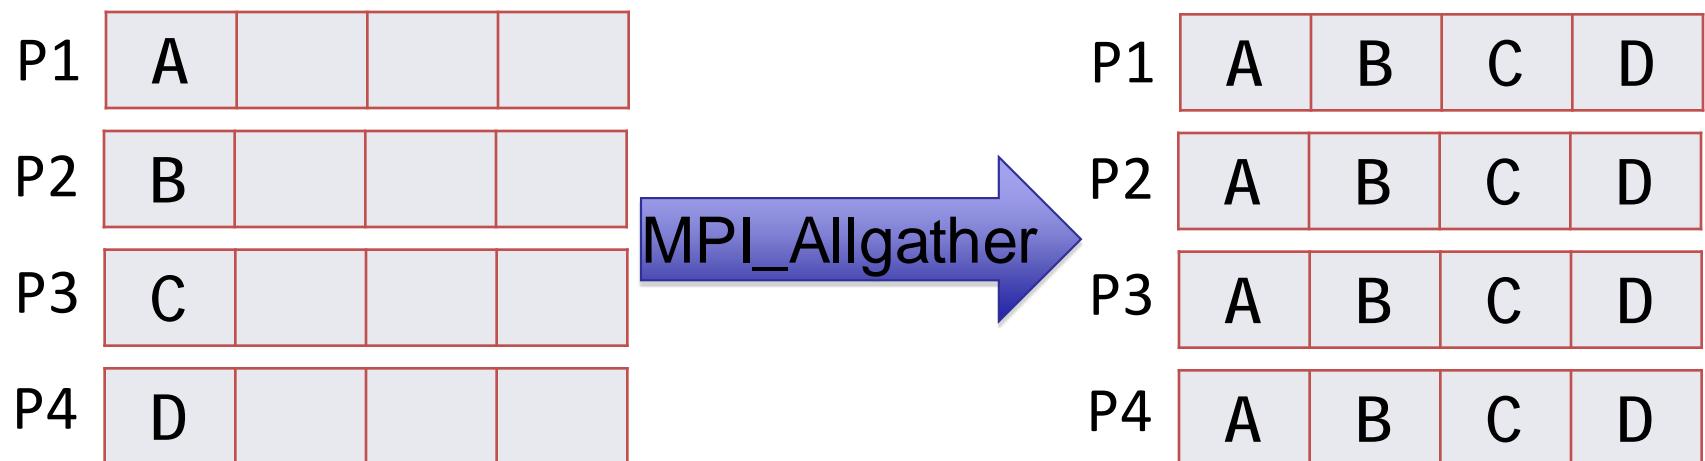
# Gather

- `int MPI_Gather(void *sendbuf, int sendcnt, MPI_Datatype sendtype, void *recvbuf, int recvcnt, MPI_Datatype recvtype, int root, MPI_Comm comm)`
  - One process (root) collects data to all the other processes in the same communicator
  - Must be called by all the processes with the same arguments



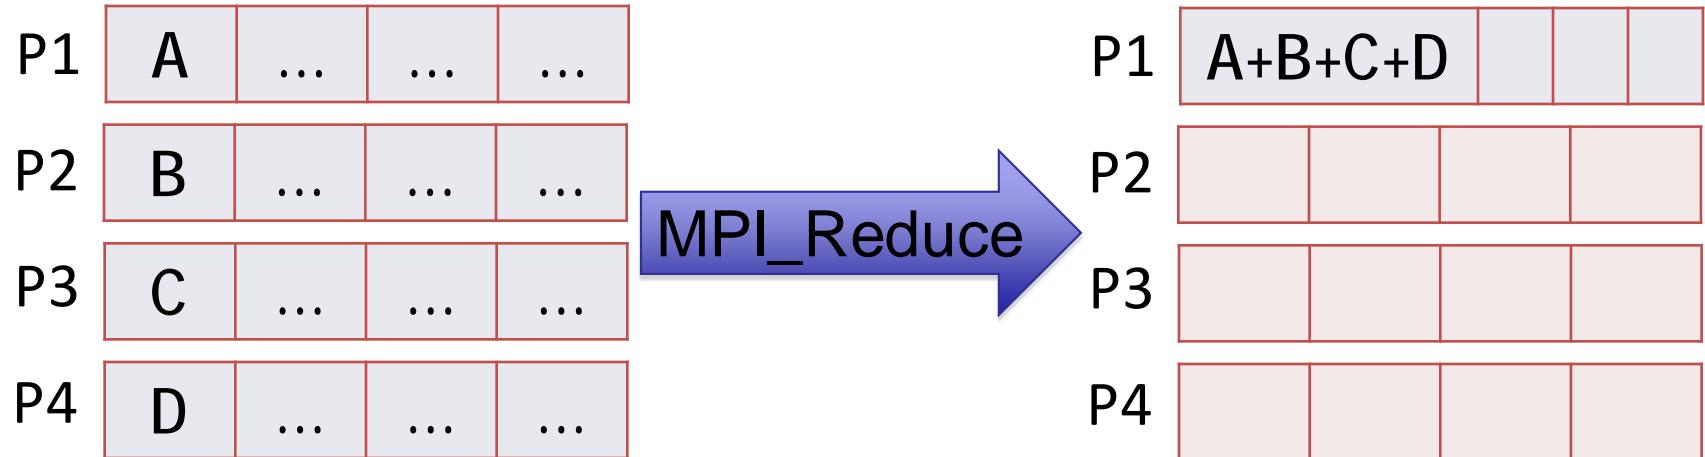
# Gather to All

- `int MPI_Allgather(void *sendbuf, int sendcnt, MPI_Datatype sendtype, void *recvbuf, int recvcnt, MPI_Datatype recvtype, MPI_Comm comm)`
  - All the processes collects data to all the other processes in the same communicator
  - Must be called by all the processes with the same arguments



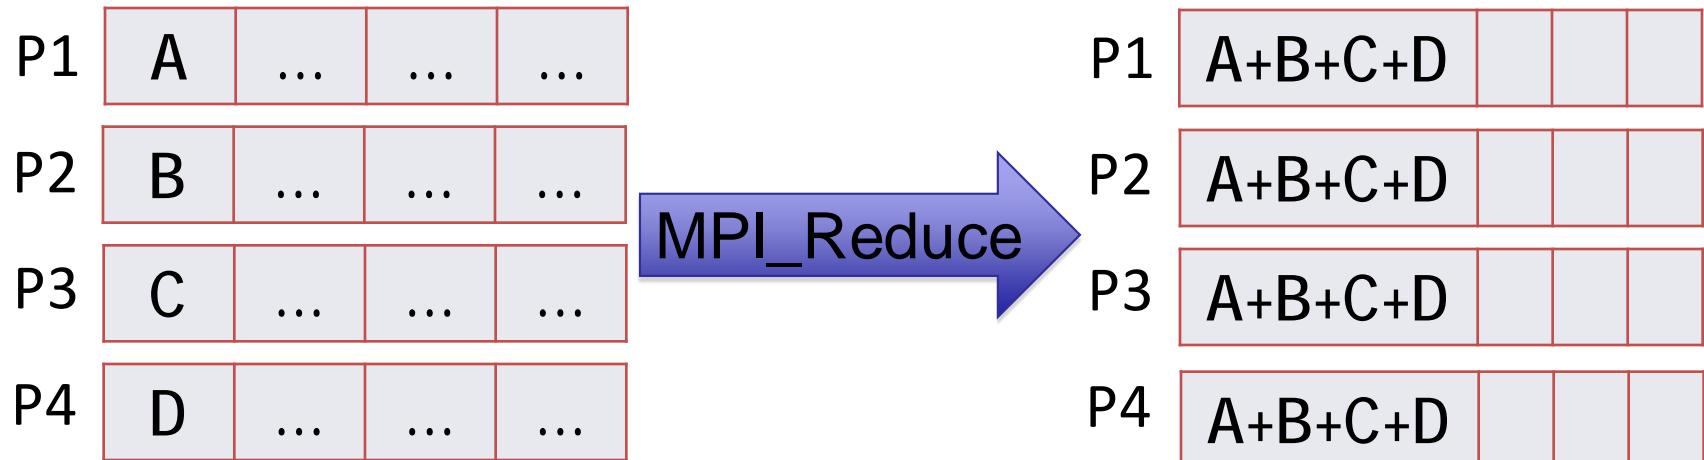
# Reduction

- `int MPI_Reduce(void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)`
  - One process (root) collects data to all the other processes in the same communicator, and performs an operation on the data
  - `MPI_SUM`, `MPI_MIN`, `MPI_MAX`, `MPI_PROD`, logical AND, OR, XOR, and a few more
  - `MPI_Op_create()`: User defined operator



# Reduction to All

- `int MPI_Allreduce(void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)`
  - All the processes collect data to all the other processes in the same communicator, and perform an operation on the data
  - `MPI_SUM`, `MPI_MIN`, `MPI_MAX`, `MPI_PROD`, logical AND, OR, XOR, and a few more
  - `MPI_Op_create()`: User defined operator



# Synchronization

- int MPI\_Barrier(MPI\_Comm comm)

```
#include "mpi.h"
#include <stdio.h>
int main(int argc, char *argv[]) {
    int rank, nprocs;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&nprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&rank);
    MPI_Barrier(MPI_COMM_WORLD);
    printf("Hello, world. I am %d of %d\n", rank, nprocs);
    MPI_Finalize();
    return 0;
}
```

# Some concepts

- The default communicator is the `MPI_COMM_WORLD`
- A process is identified by its rank in the group associated with a communicator.

# Data Types

- The data message which is sent or received is described by a triple (address, count, datatype).
- The following data types are supported by MPI:
  - Predefined data types that are corresponding to data types from the programming language.
  - Arrays.
  - Sub blocks of a matrix
  - User defined data structure.
  - A set of predefined data types

# Basic MPI types

## MPI datatype

MPI\_CHAR  
MPI\_SIGNED\_CHAR  
MPI\_UNSIGNED\_CHAR  
MPI\_SHORT  
MPI\_UNSIGNED\_SHORT  
MPI\_INT  
MPI\_UNSIGNED  
MPI\_LONG  
MPI\_UNSIGNED\_LONG  
MPI\_FLOAT  
MPI\_DOUBLE  
MPI\_LONG\_DOUBLE

## C datatype

signed char  
signed char  
unsigned char  
signed short  
unsigned short  
signed int  
unsigned int  
signed long  
unsigned long  
float  
double  
long double

# Why defining the data types during the send of a message?

Because communications take place between heterogeneous machines. Which may have different data representation and length in the memory.

# MPI blocking send

```
MPI_SEND(void *start, int  
        count, MPI_Datatype datatype, int dest,  
        int tag, MPI_Comm comm)
```

- The message buffer is described by (**start**, **count**, **datatype**).
- **dest** is the rank of the target process in the defined communicator.
- **tag** is the message identification number.

# MPI blocking receive

```
MPI_RECV(void *start, int count,  
MPI_Datatype datatype, int source, int tag,  
MPI_Comm comm, MPI_Status *status)
```

- **Source** is the rank of the sender in the communicator.
- The receiver can specify a wildcard value for source (MPI\_ANY\_SOURCE) and/or a wildcard value for tag (MPI\_ANY\_TAG), indicating that any source and/or tag are acceptable
- **Status** is used for extra information about the received message if a wildcard receive mode is used.
- If the count of the message received is less than or equal to that described by the MPI receive command, then the message is successfully received. Else it is considered as a buffer overflow error.

# MPI\_STATUS

- Status is a data structure
- In C:

```
int recvd_tag, recvd_from, recvd_count;  
MPI_Status status;  
MPI_Recv(..., MPI_ANY_SOURCE, MPI_ANY_TAG, ...,  
         &status)  
recvд_tag = status.MPI_TAG;  
recvд_from = status.MPI_SOURCE;  
MPI_Get_count(&status, datatype, &recvд_count);
```

# More info

- A receive operation may accept messages from an arbitrary sender, but a send operation must specify a unique receiver.
- Source equals destination is allowed, that is, a process can send a message to itself.

# Why MPI is simple?

- Many parallel programs can be written using just these six functions, only two of which are non-trivial;
  - MPI\_INIT
  - MPI\_FINALIZE
  - MPI\_COMM\_SIZE
  - MPI\_COMM\_RANK
  - MPI\_SEND
  - MPI\_RECV

# Non-Blocking Send and Receive

```
MPI_ISEND(buf, count, datatype, dest, tag, comm,  
          request)
```

```
MPI_IRECV(buf, count, datatype, dest, tag, comm,  
           request)
```

- **request** is a request handle which can be used to query the status of the communication or wait for its completion.

# Non-Blocking Send and Receive (Cont.)

- A non-blocking send call indicates that the system may start copying data out of the send buffer. The sender must not access any part of the send buffer after a non-blocking send operation is posted, until the complete-send returns.
- A non-blocking receive indicates that the system may start writing data into the receive buffer. The receiver must not access any part of the receive buffer after a non-blocking receive operation is posted, until the complete-receive returns.

# Non-Blocking Send and Receive (Cont.)

```
MPI_WAIT (request, status)
```

```
MPI_TEST (request, flag, status)
```

- The MPI\_WAIT will block your program until the non-blocking send/receive with the desired request is done.
- The MPI\_TEST is simply queried to see if the communication has completed and the result of the query (TRUE or FALSE) is returned immediately in flag.

# Deadlocks in blocking operations

- What happens with

Process 0

Send(1)

Recv(1)

Process 1

Send(0)

Recv(0)

- Send a large message from process 0 to process 1
  - If there is insufficient storage at the destination, the send must wait for the user to provide the memory space(through a receive)
- This is called “unsafe” because it depends on the availability of system buffers.

# Some solutions to the “unsafe” problem

- Order the operations more carefully

Process 0

Send(1)

Recv(1)

Process 1

Recv(0)

Send(0)

Use non-blocking operations:

Process 0

ISend(1)

IRecv(1)

Waitall

Process 1

ISend(0)

IRecv(0)

Waitall

# Introduction to collective operations in MPI

- o Collective operations are called by all processes in a communicator
- o MPI\_Bcast distributes data from one process(the root) to all others in a communicator.

Syntax:

```
MPI_Bcast(void *message, int count, MPI_Datatype  
datatype, int root, MPI_Comm comm)
```

- o MPI\_Reduce combines data from all processes in communicator or and returns it to one process

Syntax:

```
MPI_Reduce(void *message, void *recvbuf, int count,  
MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm  
comm)
```

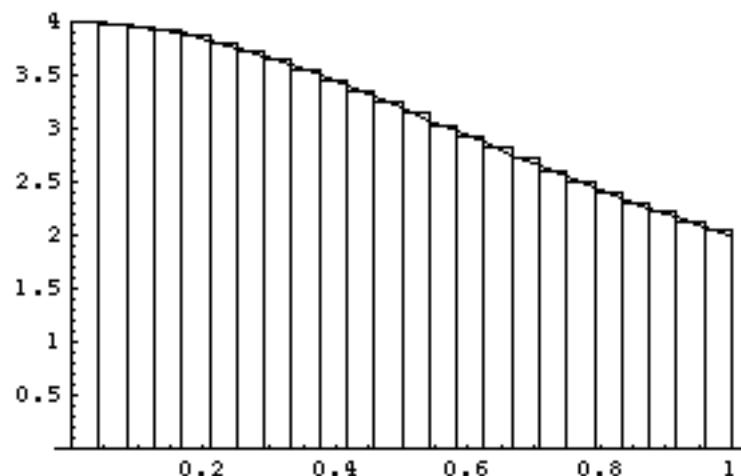
- o In many numerical algorithm, send/receive can be replaced by Bcast/Reduce, improving both simplicity and efficiency.

# Collective Operations

MPI\_MAX, MPI\_MIN, MPI\_SUM, MPI\_PROD, MPI LAND,  
MPI\_BAND, MPI\_LOR, MPI\_BOR, MPI\_LXOR, MPI\_BXOR,  
MPI\_MAXLOC, MPI\_MINLOC

# Example: Compute PI (0)

$$\pi = \int_0^1 \frac{4}{1+x^2} dx$$



# Example: Compute PI (1)

```
#include "mpi.h"
#include <math.h>

int main(int argc, char *argv[])
{
    int done = 0, n, myid, numprocs, I, rc;
    double PI25DT = 3.141592653589793238462643;
    double mypi, pi, h, sum, x, a;
    MPI_INIT(&argc, &argv);
    MPI_COMM_SIZE(MPI_COMM_WORLD, &numprocs);
    MPI_COMM_RANK(MPI_COMM_WORLD, &myid);
    while (!done)
    {
        if (myid == 0)
        {
            printf("Enter the number of intervals: (0 quits) ");
            scanf("%d", &n);
        }
        MPI_BCAST(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
        if (n == 0)
    }
```

# Example: Compute PI (2)

```
h = 1.0 / (double)n;
sum = 0.0;
for (i = myid + 1; i <= n; i += numprocs)
{
    x = h * ((double)i - 0.5);
    sum += 4.0 / (1.0 + x * x);
}
mypi = h * sum;
MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0,
MPI_COMM_WORLD);

if (myid == 0) printf("pi is approximately %.16f, Error
is %.16f\n", pi, fabs(pi - PI25DT));

MPI_Finalize();
return 0;
}
```

# When to use MPI

- Portability and Performance
- Irregular data structure
- Building tools for others
- Need to manage memory on a per processor basis

# Examples

- Simple

```
#include "main.h"

using namespace std;

int main(int argc, char* argv[]){
    int mytid, numprocs;
    MPI_Status status;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&mytid);

    char name[100];
    gethostname(name, sizeof(name));

    if(mytid>0){
        funkywork();
    }

    cout << "Hello, " << mytid << " and " << name << " say hi in a C++ statement\n";

    MPI_Finalize();

}
```

# • Serial v/s MPI parallel

```
#include <stdio.h>

#define NRA 5000          /* number of rows in matrix A */
#define NCA 1000          /* number of columns in matrix A */
#define NCB 100           /* number of columns in matrix B */

main()
{
    int i,j,k;          /* misc */
    double a[NRA][NCA],   /* matrix A to be multiplied */
           b[NCA][NCB],   /* matrix B to be multiplied */
           c[NRA][NCB];   /* result matrix C */

    /* Initialize A, B, and C matrices */
    for (i=0;i<NRA;i++)
        for (j=0;j<NCA;j++)
            a[i][j]=i+j;
    for (i=0;i<NCA;i++)
        for (j=0;j<NCB;j++)
            b[i][j]=i*j;
    for(i=0;i<NRA;i++)
        for(j=0;j<NCB;j++)
            c[i][j]=0.0;

    /* Perform matrix multiply */
    for(i=0;i<NRA;i++)
        for(j=0;j<NCB;j++)
            for(k=0;k<NCA;k++)
                c[i][j]+=a[i][k] * b[k][j];

    /* Okay, it's a trivial program */
    printf("Here is the result matrix\n");
    for (i=0;i<NRA;i++)
    {
        printf("\n");
        for (j=0;j<NCB;j++)
            printf("%6.2f ",c[i][j]);
    }
    printf("\n");
}
```

```
#include "mpi.h"
#include <stdio.h>
#define NRA 5000          /* number of rows in matrix A */
#define NCA 1000          /* number of columns in matrix A */
#define NCB 100           /* number of columns in matrix B */
#define MASTER 0           /* taskid of first task */
#define FROM_MASTER 1     /* setting a message type */
#define FROM_WORKER 2     /* setting a message type */

int main(argc,argv)
int argc;
char *argv[];
{
    int numtasks,          /* number of tasks in partition */
        taskid,           /* a task identifier */
        numworkers,         /* number of worker tasks */
        source,             /* task id of message source */
        dest,               /* task id of message destination */
        mtype,              /* message type */
        rows,               /* rows of matrix A sent to each worker */
        averow, extra, offset, /* used to determine rows sent to each
                               worker */
        i, j, k, rc;         /* misc */
    double a[NRA][NCA],   /* matrix A to be multiplied */
           b[NCA][NCB],   /* matrix B to be multiplied */
           c[NRA][NCB];   /* result matrix C */
    MPI_Status status;

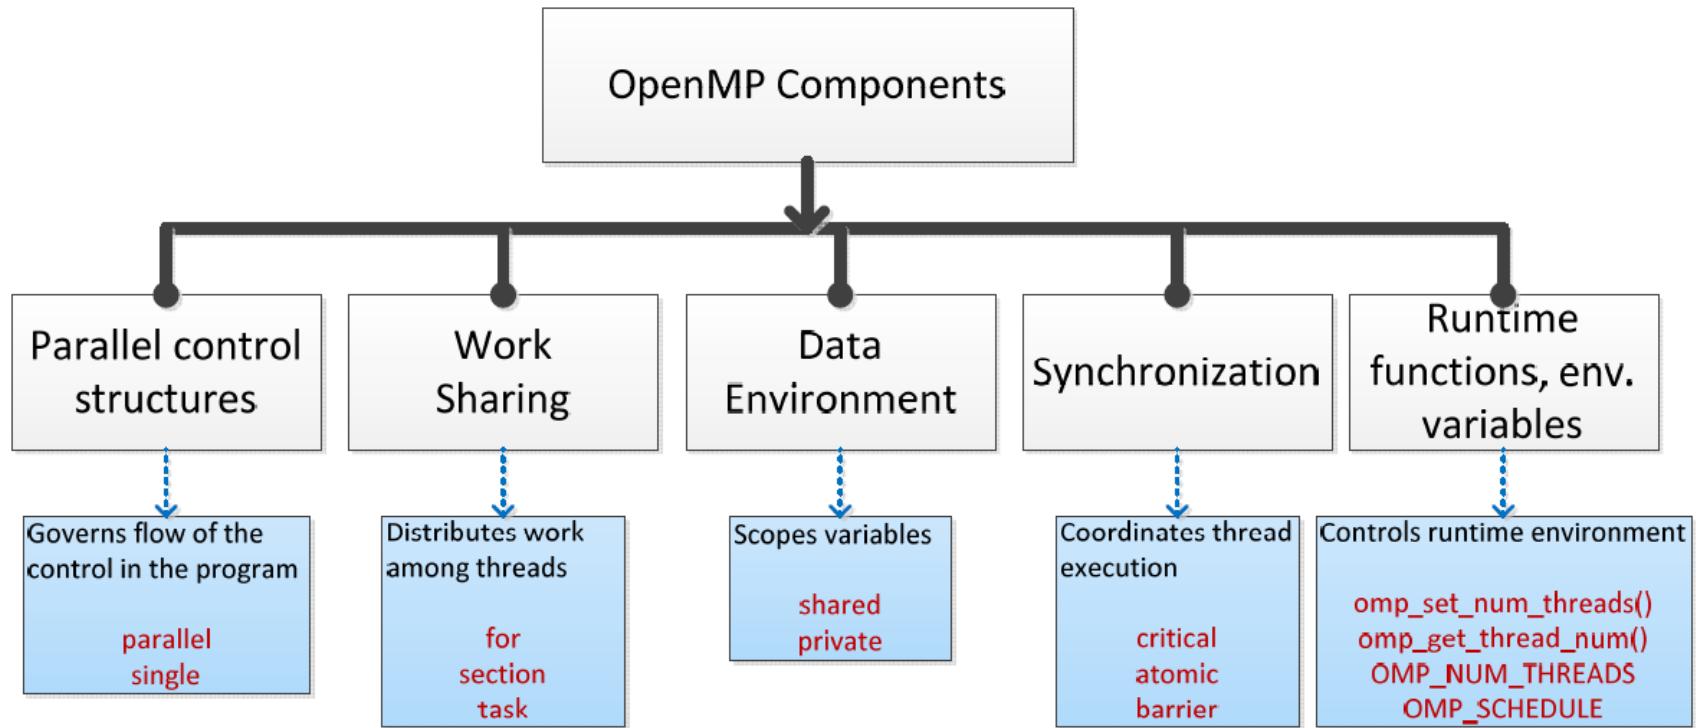
    rc = MPI_Init(&argc,&argv);
    rc |= MPI_Comm_size(MPI_COMM_WORLD,&numtasks);
    rc |= MPI_Comm_rank(MPI_COMM_WORLD,&taskid);
    if (rc != 0)
        printf ("error initializing MPI and obtaining task ID
information\n");
    else
        printf ("task ID = %d\n", taskid);
    numworkers = numtasks-1;
    /***** master task *****/
    *****/
}
```

LETS RUN THIS ON DISCOVERY CLUSTER

# OpenMP

- v2.5, v3.0 and v3.1 standards
- gcc 4.4 and higher implements v3.0, and 4.7 v3.1
- Syntax:**#pragma omp directive-name [clause[ ,] clause]... ] new-line**
- Execution Model: OpenMP API uses the fork-join model of parallel execution
- Memory Model: OpenMP has a relaxed-consistency, shared-memory model
- Each thread has temporary view of memory and thread private memory
- OpenMP flush - enforces consistency between temporary view and memory

[http://gcc.gnu.org/onlinedocs/libgomp/index.html#toc\\_Runtime-Library-Routines](http://gcc.gnu.org/onlinedocs/libgomp/index.html#toc_Runtime-Library-Routines)



- Application Programmer Interface (API) is combination of
  - Directives
    - Example: `#pragma omp task`
  - Runtime library routines
    - Example: `int omp_get_thread_num(void)`
  - Environment variables
    - Example: `setenv OMP_SCHEDULE "guided, 4"`

- Directives (or Pragmas) used to
  - Express/Define parallelism (flow control)
    - Example: `#pragma omp parallel for`
  - Specify data sharing among threads (communication)
    - Example: `#pragma omp parallel for private(x,y)`
  - Synchronization (coordination or interaction)
    - Example: `#pragma omp barrier`

# Summary of Run-Time Library OpenMP Routines

1. `omp_set_num_threads`
2. `omp_get_num_threads`
3. `omp_get_max_threads`
4. `omp_get_thread_num`
5. `omp_get_thread_limit`
6. `omp_get_num_procs`
7. `omp_in_parallel`
8. `omp_set_dynamic`
9. `omp_get_dynamic`
10. `omp_set_nested`
11. `omp_get_nested`
12. `omp_set_schedule`
13. `omp_get_schedule`
14. `omp_set_max_active_levels`
15. `omp_get_max_active_levels`
16. `omp_get_level`
17. `omp_get_ancestor_thread_num`
18. `omp_get_team_size`
19. `omp_get_active_level`
20. `omp_init_lock`
21. `omp_destroy_lock`
22. `omp_set_lock`
23. `omp_unset_lock`
24. `omp_test_lock`
25. `omp_init_nest_lock`
26. `omp_destroy_nest_lock`
27. `omp_set_nest_lock`
28. `omp_unset_nest_lock`
29. `omp_test_nest_lock`
30. `omp_get_wtime`
31. `omp_get_wtick`

## OMP\_SCHEDULE

- Example: `setenv OMP_SCHEDULE "guided, 4"`

## OMP\_NUM\_THREADS

- Sets the maximum number of threads to use during execution.
- Example: `setenv OMP_NUM_THREADS 8`

## OMP\_DYNAMIC

- Enables or disables dynamic adjustment of the number of threads available for execution of parallel regions. Valid values are TRUE or FALSE
- Example: `setenv OMP_DYNAMIC TRUE`

## OMP\_NESTED

- Enables or disables nested parallelism. Valid values are TRUE or FALSE
- Example: `setenv OMP_NESTED TRUE`

## OMP\_STACKSIZE

- Controls the size [in KB] of the stack for created (non-Master) threads.

## OMP\_WAIT\_POLICY

- Provides hint to an OpenMP implementation about desired behavior of waiting threads

## OMP\_MAX\_ACTIVE\_LEVELS

- Controls the maximum number of nested active parallel regions. The value of this environment variable must be a non-negative integer. Example:
- `setenv OMP_MAX_ACTIVE_LEVELS 2`

## OMP\_THREAD\_LIMIT

- Sets the number of OpenMP threads to use for the whole OpenMP program Example:
- `setenv OMP_THREAD_LIMIT 8`

## Serial regions by default, annotate to create parallel regions

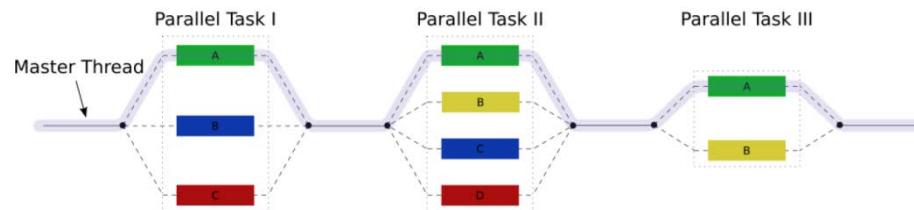
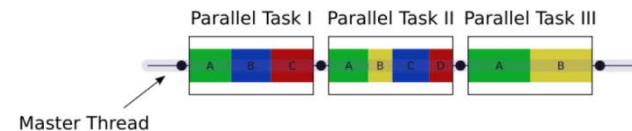
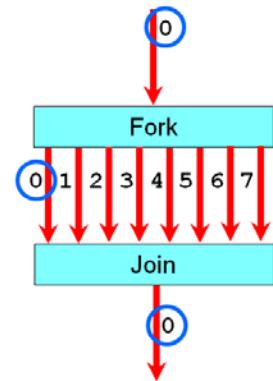
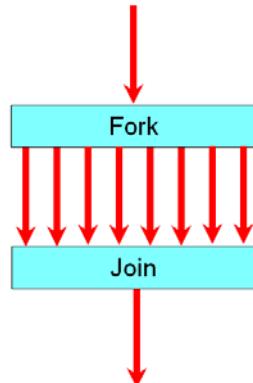
- Generic parallel regions
- Parallelized loops
- Sectioned parallel regions

## Thread-like Fork/Join model

- Arbitrary number of *logical* thread creation/destruction events

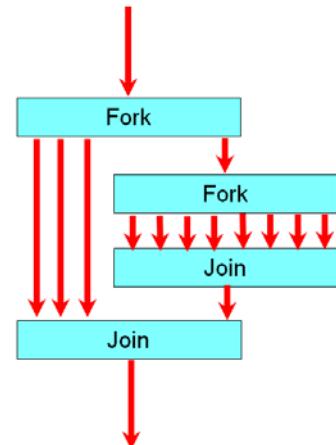
### Master Thread

- Thread with ID=0
- Only thread that exists in sequential regions
- Depending on implementation, may have special purpose inside parallel regions
- Some special directives affect only the master thread (like `master`)



## Fork/Join can be nested

- Nesting complication handled “automagically” at compile-time
- Independent of the number of threads actually running

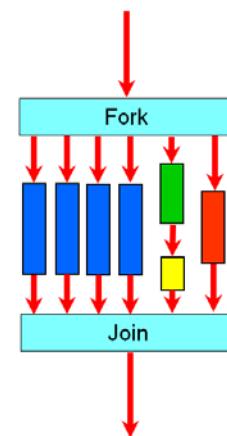


### Data parallelism

- Threads perform similar functions, guided by thread identifier

### Control parallelism

- Threads perform differing functions
  - » One thread for I/O, one for computation, etc...



# OpenMP Loop Scheduling

**schedule clause determines how loop iterations are divided among the thread team**

- static([chunk]) divides iterations statically between threads
  - » Each thread receives [chunk] iterations, rounding as necessary to account for all iterations
  - » Default [chunk] is  $\text{ceil}(\# \text{ iterations} / \# \text{ threads})$
- dynamic([chunk]) allocates [chunk] iterations per thread, allocating an additional [chunk] iterations when a thread finishes
  - » Forms a logical work queue, consisting of all loop iterations
  - » Default [chunk] is 1
- guided([chunk]) allocates dynamically, but [chunk] is exponentially reduced with each allocation

```
#pragma omp parallel for \
schedule(static)
for( i=0; i<16; i++ )
{
    dolteration(i);
}
```

```
// Static Scheduling
int chunk = 16/T;
int base = tid * chunk;
int bound = (tid+1)*chunk;
for( i=base; i<bound; i++ )
{
    dolteration(i);
}
```

```
Barrier();
```

```
#pragma omp parallel for \
schedule(dynamic)
for( i=0; i<16; i++ )
{
    dolteration(i);
}
```

```
// Dynamic Scheduling
int current_i;
while( workLeftToDo() )
{
    current_i = getNextIter();
    dolteration(i);
}
Barrier();
```

# OpenMP Data Sharing

**Parallel programs often employ two types of data**

- Shared data, visible to all threads, similarly named
- Private data, visible to a single thread (often stack-allocated)

**PThreads:**

- Global-scoped variables are shared
- Stack-allocated variables are private

**OpenMP:**

- shared variables are shared
- private variables are private

# OpenMP Synchronization

- OpenMP Critical Sections
  - » Named or unnamed
  - » No *explicit* locks
- Barrier directives
- Explicit Lock functions
  - » When all else fails – may require flush directive
- Single-thread regions *within* parallel regions
  - » master, single directives

```
#pragma omp critical
{
    /* Critical code here */
}

#pragma omp barrier

omp_set_lock( lock l );
/* Code goes here */

omp_unset_lock( lock l );
#pragma omp single
{
    /* Only executed once */
}
```

# OpenMP Program to Compute PI

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>
int main (int argc, char *argv[])
{
    int nthreads, tid;
    int i, INTERVALS;
    double n_1, x, pi = 0.0;

INTERVALS=128000;
/* Fork a team of threads giving them their own
   copies of variables */
#pragma omp parallel private(nthreads, tid)
{
    /* Obtain thread number */
    tid = omp_get_thread_num();
    printf("Hello from thread = %d\n", tid);

    /* Only master thread does this */
    if (tid == 0)
    {
        nthreads = omp_get_num_threads();
        printf("Number of threads = %d\n", nthreads);
    }

    } /* All threads join master thread and disband */

    n_1 = 1.0 / (double)INTERVALS;

    /* Parallel loop with reduction for calculating PI */
#pragma omp parallel for private(i,x) shared
(n_1,INTERVALS) reduction(+:pi)
    for (i = 0; i < INTERVALS; i++)
    {
        x = n_1 * ((double)i - 0.5);
        pi += 4.0 / (1.0 + x * x);
    }
    pi *= n_1;
    printf ("Pi = %.12f\n", pi);
}
```

LET'S RUN AN EXAMPLE OF MATRIX  
MULTIPLICATION IN SERIAL AND WITH  
OpenMP PARALLELIZATION ON DISCOVERY  
CLUSTER

# Hybrid OpenMP + MPI

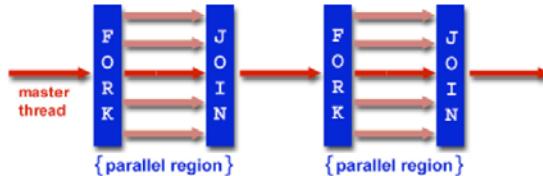
```
#include<mpi.h> /* MPI Library */  
#include<omp.h> /* OpenMP Library */  
#include<stdio.h> /* printf() */  
#include<stdlib.h> /* EXIT_SUCCESS */  
  
int main (int argc, char *argv[]) {  
  
    /* Parameters of MPI. */  
    int M_N; /* number of MPI ranks */  
    int M_ID; /* MPI rank ID */  
    int rtn_val; /* return value */  
    char name[128]; /* MPI_MAX_PROCESSOR_NAME == 128 */  
    int namelen;  
  
    /* Parameters of OpenMP. */  
    int O_P; /* number of OpenMP processors */  
    int O_T; /* number of OpenMP threads */  
    int O_ID; /* OpenMP threadID */  
  
    /* Initialize MPI. */  
    /* Construct the default communicator MPI_COMM_WORLD. */  
    rtn_val = MPI_Init(&argc,&argv);  
  
    /* Get a few MPI parameters. */  
    rtn_val = MPI_Comm_size(MPI_COMM_WORLD,&M_N); /* get number of MPI ranks */  
    rtn_val = MPI_Comm_rank(MPI_COMM_WORLD,&M_ID); /* get MPI rankID */  
    MPI_Get_processor_name(name,&namelen);  
    printf("name:%s M_ID:%d M_N:%d\n",name,M_ID,M_N);  
  
    /* Get a few OpenMP parameters. */  
    O_P = omp_get_num_procs(); /* get number of OpenMP processors */  
    O_T = omp_get_num_threads(); /* get number of OpenMP threads */  
    O_ID=omp_get_thread_num(); /* get OpenMP threadID */  
    printf("name:%s M_ID:%d O_ID:%d O_P:%d O_T:%d\n",name,M_ID,O_ID,O_P,O_T);  
  
    /* PARALLEL REGION */  
    /* Thread IDs range from 0 through omp_get_num_threads()-1. */  
    /* We execute identical code in all threads (data parallelization). */  
    #pragma omp parallel private(O_ID)  
    {  
        O_ID=omp_get_thread_num(); /* get OpenMP threadID */  
        MPI_Get_processor_name(name,&namelen);  
        printf("parallel region: name:%s M_ID=%d O_ID=%d\n",name,M_ID,O_ID);  
    }  
  
    /* Terminate MPI. */  
    rtn_val = MPI_Finalize();  
  
    /* Exit master thread. */  
    printf("name:%s M_ID:%d O_ID:%d Exits\n",name,M_ID,O_ID);  
    return EXIT_SUCCESS;  
}
```

# OpenMP Summary

Shared memory, thread-based parallelism

Explicit parallelism (relies on you specifying parallel regions)

Fork/join model



Industry-standard shared memory programming model

- First version released in 1997

OpenMP Architecture Review Board (ARB) determines updates to standard

- The final specification of Version 3.1 released in July of 2011 (minor update)

OpenMP provides small yet versatile programming model

- This model serves as the inspiration for the OpenACC effort to standardizing approaches that can factor in the presence of a GPU accelerator

Not at all intrusive, very straightforward to parallelize existing code

- Good efficiency gains achieved by using parallel regions in an existing code

Work-sharing constructs: `for`, `section`, `task` enable parallelization of computationally intensive portions of program

Parallelize small parts of application, one at a time (beginning with most time-critical parts)

Can implement complex algorithms

Code size grows only modestly

Expression of parallelism flows clearly, code is easy to read

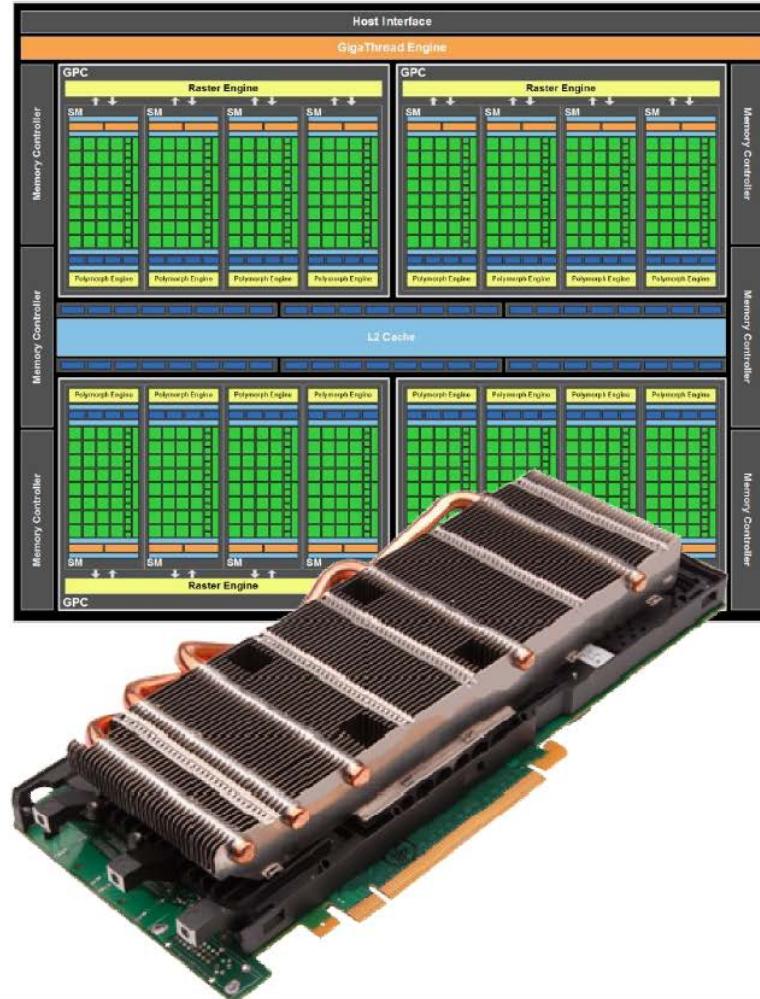
OpenMP threads are heavy

- Very good for handling parallel tasks
- Not particularly remarkable at handling fine grain data parallelism (vector architectures excel here)

# GPGPU

## NVIDIA GPUs

- Supports CUDA and OpenCL
- Fermi (Tesla version)
  - Up to 512 cores
  - DP **0.5 Tflop/s**
  - **3-6 GB of memory**
  - **Caches included**
    - L1 per multiprocessor
    - L2: Shared
- Kepler in 2012
- Maxwell in 2014

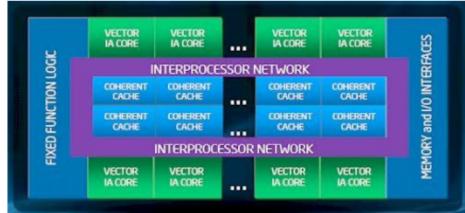


NVIDIA CUDA 6.5 SDK

## Other accelerators

### • Intel

- Intel MIC (Many Integrated Core)
- More than 50 X86 vector cores
- OpenMP, OpenCL, Intel parallel building blocks...
- First product (Knights Corner) in 2012



### • AMD

- Line of GPUs supporting OpenCL
- Not so widely used in HPC at the moment

| Generation | Model                   | Video card equivalent | GPU Core        | Threads max.           | Core                |             | Memory            |       |                 |              | Raw processing power<br>(Floating-Point Operations per Second) |                     | Peak TDP (watts)         | Others |                   |
|------------|-------------------------|-----------------------|-----------------|------------------------|---------------------|-------------|-------------------|-------|-----------------|--------------|----------------------------------------------------------------|---------------------|--------------------------|--------|-------------------|
|            |                         |                       |                 |                        | SPUs <sup>NB1</sup> | Clock (MHz) | Bandwidth (GiB/s) | Type  | Bus width (bit) | Amount (MiB) | Clock (MHz)                                                    | FP32 GFLOPs         | FP64 GFLOPs              |        |                   |
| 1st NB2    | 580 <sup>[7][10]</sup>  | Radeon X1900 XTX      | R580            | 512                    | 48                  | 600         | 83.2              | GDDR3 | 256             | 1024         | 650                                                            | 375 <sup>[11]</sup> | N/A                      | ≤165   |                   |
| 2nd NB2    | 9170 <sup>[9][12]</sup> | Radeon HD 3870        | RV670           | ?                      | 64 (320)            | 800         | 51.2              | GDDR3 | 256             | 2048         | 800                                                            | 512                 | 102.4 <sup>NB3[13]</sup> | ≤105   |                   |
| 3rd NB2    | 9250 <sup>[14]</sup>    | Radeon HD 4850        | RV770           | 16,384 <sup>[15]</sup> | 160 (800)           | 625         | 63.5              | GDDR3 | 256             | 1024         | 993                                                            | 1000                | 200 <sup>NB3</sup>       | ≤150   |                   |
|            | 9270 <sup>[16]</sup>    | Radeon HD 4870        |                 |                        | 750                 | 108.8       | GDDR5             | 256   | 2048            | 850          | 1200                                                           | 240 <sup>NB3</sup>  | <160                     |        |                   |
| 4th NB2    | 9350                    | Radeon HD 5850        | Cypress (RV870) | 31,744 <sup>[17]</sup> | 288 (1440)          | 700         | 128               | GDDR5 | 256             | 2048         | 1000                                                           | 2016                | 403.2                    | 150    | codenamed Kestrel |
|            | 9370                    | Radeon HD 5870        |                 |                        | 320 (1600)          | 825         | 147.2             |       | 256             | 4096         | 1150                                                           | 2640                | 528                      | 225    | codenamed Osprey  |

AMD stream processing lineup - AMD Stream SDK was replaced by AMD APP SDK, available for Microsoft Windows and Linux, 32-bit and 64-bit. APP stands for "Accelerated Parallel Processing"

## Intel® Xeon Phi™ Product Family Specifications

| PRODUCT NUMBER | FORM FACTOR & THERMAL SOLUTION <sup>4</sup> | BOARD TDP (WATTS) | NUMBER OF CORES | FREQUENCY (GHz) | PEAK DOUBLE PRECISION PERFORMANCE (GFLOP) | PEAK MEMORY BANDWIDTH (GB/s) | MEMORY CAPACITY (GB) | INTEL® TURBO BOOST TECHNOLOGY  |
|----------------|---------------------------------------------|-------------------|-----------------|-----------------|-------------------------------------------|------------------------------|----------------------|--------------------------------|
| 3120P          | PCIe, Passive                               | 300               | 57              | 1.1             | 1003                                      | 240                          | 6                    | N/A                            |
| 3120A          | PCIe, Active                                | 300               | 57              | 1.1             | 1003                                      | 240                          | 6                    | N/A                            |
| 5110P          | PCIe, Passive                               | 225               | 60              | 1.053           | 1011                                      | 320                          | 8                    | N/A                            |
| 5120D          | Dense form factor, None                     | 245               | 60              | 1.053           | 1011                                      | 352                          | 8                    | N/A                            |
| 7110P          | PCIe, Passive                               | 300               | 61              | 1.238           | 1208                                      | 352                          | 16                   |                                |
| 7120X          | PCIe, None                                  | 300               | 61              | 1.238           | 1208                                      | 352                          | 16                   | Peak turbo frequency: 1.33 GHz |

## Intel Compilers and Intel Parallel Studio

# How to use GPUs

1. Use existing GPU software
2. Use numerical libraries for GPUs
3. Program GPU code with directives
4. Program native GPU code



## Directive based GPU code

- Two main products
  - PGI accelerator
  - HMPP (CAPS enterprise)
- Normal C or Fortran code with directives to guide compiler in creating a GPU version
- Backends supporting CUDA, OpenCL and even normal CPUs

```
//HMPP codelet
#pragma hmpp label1 codelet, args[B].io=out, args [C].io=inout,
target=CUDA:CAL/IL
void myFunc(int n, int A[n], int B[n], int C[n]) {
    for(int i=0 ; i<n ; i++) {
        B[i] = A[i] * A[i]; C[i] = C[i] * A[i];
    }
}
```

```
//HMPP callsite
#pragma hmpp label1 callsite
myFunc(n, A, B, C);
```

## Native GPU code

- CUDA, CUDA-Fortran (PGI), OpenCL
- Cons
  - Requires most time
- Pros
  - Good control & performance
  - Can be combined with library & directive approaches

## Use existing GPU software

- HOOMD, NAMD, GROMACS, GPU-HMMER, TeraChem, Matlab (jacket etc)...
- Pros
  - No implementation headaches for end users
- Cons
  - Existing applications do not cover all science areas
  - Often include limited number of algorithms/models
  - For many applications the GPU version is still immature

## Use GPU libraries

- CUBLAS, MAGMA (Lapack for GPU), ...
- Pros
  - Easy to implement in your code
  - Algorithms in libraries efficient
- Cons
  - Speedup limited by Amdahls law & transfer bottleneck

# CUDA vs OpenCL

- **CUDA Nvidia specific**
- **OpenCL is a standard adopted by all major players**
  - Writing OpenCL code providing good performance on all platforms (Nvidia, AMD, etc.) is difficult
- **On Nvidia HW**
  - CUDA is faster (at the moment)
  - According to Nvidia
    - CUDA will evolve faster than OpenCL
    - Remains Nvidias main programming language

# CUDA

---

- **Compute Unified Device Architecture**
- CUDA C is a C/C++ language extension for GPU programming
  - PGI has developed similar Fortran 2003 extension
- CUDA API is the most up-to-date GPGPU programming interface for NVIDIA GPUs
  - Two APIs: runtime and driver

# CUDA C

## Qualifiers

global, device, shared,  
local, constant, ...

## Built-in variables

threadIdx, blockIdx, ...

## Intrinsics

`__syncthreads,`  
`__fmul_rn,` ...

## Runtime API

memory, device,  
execution management

## Kernel launch

```
__device__ float array[128];
__global__ void kern(float *data) {
    __shared__ float buffer[32];

    ...
    buffer[threadIdx.x] = data[i];
    ...

    __syncthreads;

    ...
}

float *d_data;
cudaMalloc((void**)&d_data, bytes);

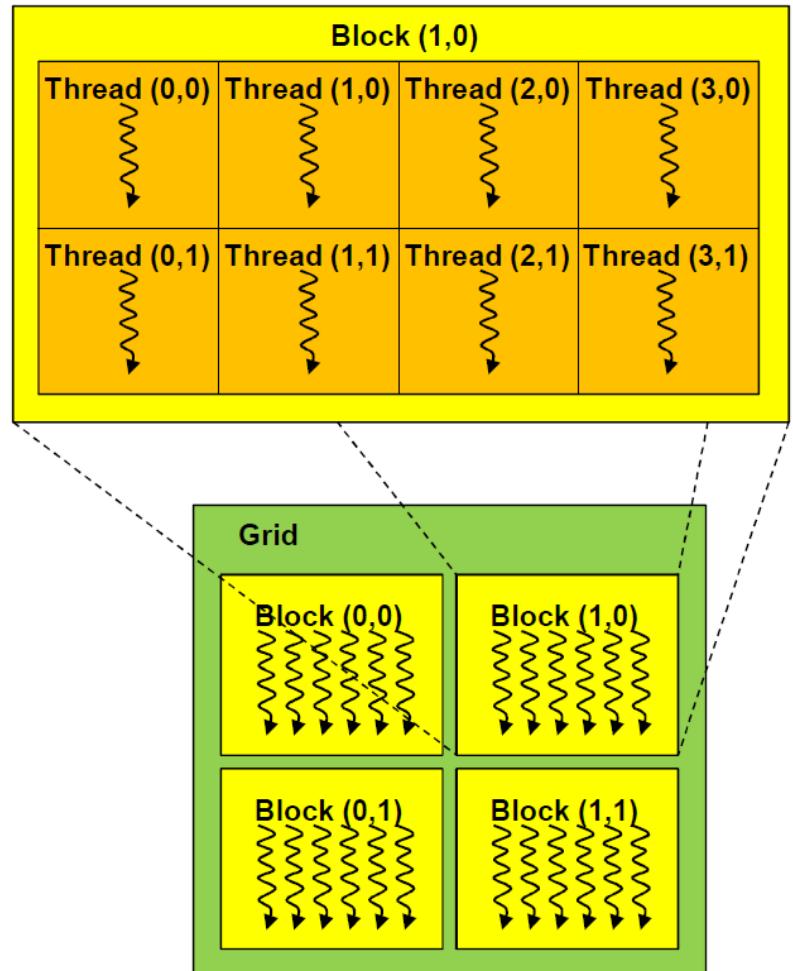
kern<<<1024, 128>>>(d_data);
```

# CUDA PROGRAMMING MODEL

- GPU accelerator is called **device**, CPU is **host**
- GPU code (**kernel**) is launched and executed on the device by several threads
- Threads are grouped into thread blocks
- Program code is written from a single thread's point of view
  - Each thread can diverge and execute a unique code path (can cause performance issues)

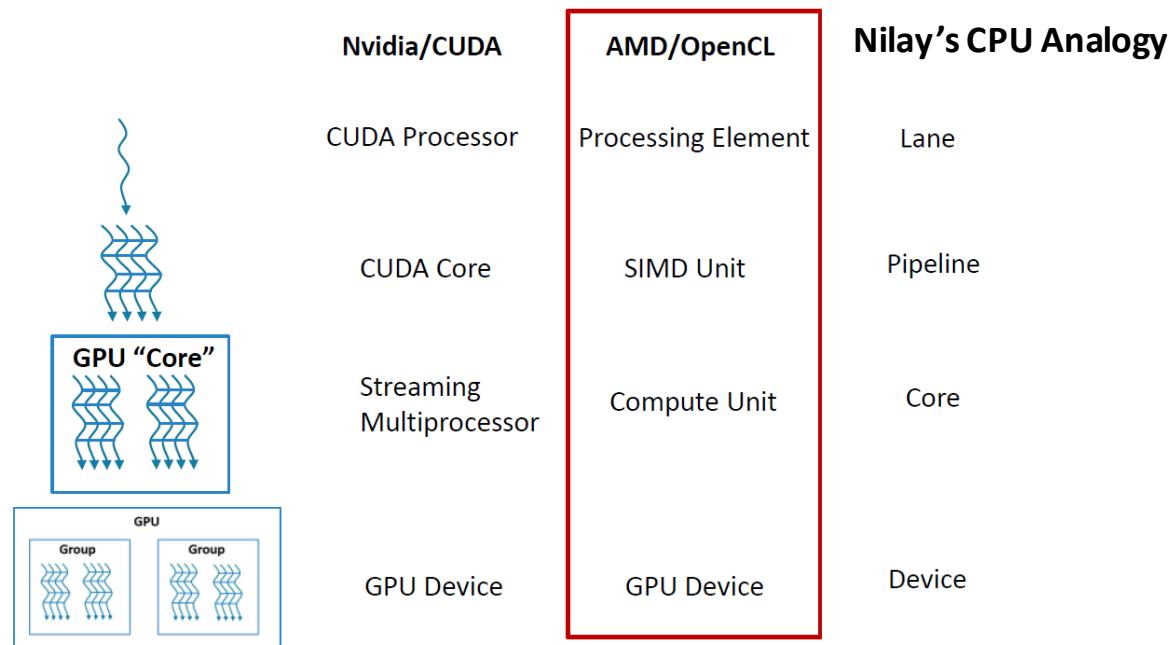
# THREAD HIERARCHY

- Threads:
  - 3D IDs, unique in `block`
- Blocks:
  - 3D\* IDs, unique in `grid`
- Dimensions are set at kernel launch
- Built-in variables for device code:
  - `threadIdx`, `blockIdx`
  - `blockDim`, `gridDim`



# Hardware Implementation, SIMT Architecture

- Maximum number of threads in a block depends on the compute capability (1024 on Fermi)
- GPU multiprocessor creates, manages, schedules and executes threads in **warps** of 32\*
- Warp executes one common instruction at a time
  - Threads are allowed to branch, but branches are serialized
- Context switch is fast, scheduler selects warps that are ready to execute → can hide latencies



# Thread Branching

Thread 0 Threads 1-31



Program

```
int tid = threadIdx.x  
  
if (tid == 0) {++var1;}  
  
else {var1 = var1 + 2;}  
  
var2 = 3 * var1;
```

Host and device have separate memories

Host manages the GPU memory

Usually one has to

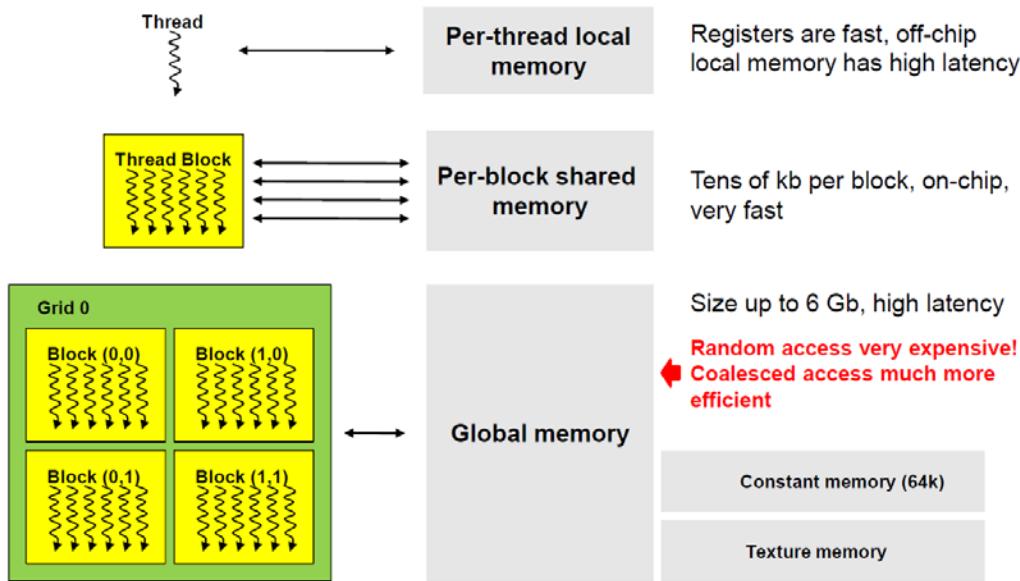
1. Copy (explicitly) data from host to the device
2. Execute the GPU kernel
3. Copy (explicitly) the results back to the host

Data copies between host and device use the PCI bus with very limited bandwidth → **minimize the transfers!**

```
int main(void) {
```

```
    float *A = (float *) malloc(N*sizeof(float));  
    float *d_A;  
    cudaMalloc((void**)&d_A, N*sizeof(float));  
    cudaMemcpy(d_A, A, N*sizeof(float), cudaMemcpyHostToDevice);  
    ...  
    float A0 = d_A[0];  
    ...  
    cudaMemcpy(A, d_A, N*sizeof(float), cudaMemcpyDeviceToHost);  
    cudaFree(d_A);  
    free(A);  
    return 0;
```

Can not dereference device pointers in host code!



# Device Code

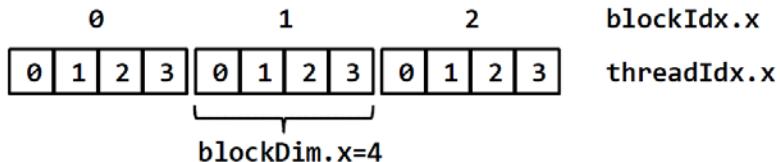
- C functions with restrictions:
  - Can only dereference pointers to device memory
  - No static variables, no recursion
  - No variable number of arguments
- Functions must be declared with a qualifier
  - `__global__`: Kernel, called from CPU
    - Cannot be called from GPU
    - Must return void
  - `__device__`: Called from `__device__` and `__global__` funcs
    - Can not be called from CPU
  - `__host__`: Can only be called by CPU
    - Can be combined with `__device__` qualifier

# Kernel Example

```
__global__ void kern(int *A) {  
    int idx = blockIdx.x * blockDim.x + threadIdx.x;  
    A[idx] = idx;  
}  
Result: A = {0,1,2,3,4,5,6,7,8,9,10,11}
```

3 blocks with 4 threads in each

```
void main() {  
    // Allocate memories, copy values  
    dim3 grid, block;  
    block.x = 4;  
    grid.x = 3;  
    kern<<<grid, block>>>(d_A);  
    // Copy results back  
}
```



```
__global__ void kern(int *A) {  
    int idx = blockIdx.x * blockDim.x + threadIdx.x;  
    A[idx] = blockIdx.x;  
}  
Result: A = {0,0,0,0,1,1,1,1,2,2,2,2}
```

```
__global__ void kern(int *A) {  
    int idx = blockIdx.x * blockDim.x + threadIdx.x;  
    A[idx] = threadIdx.x;  
}  
Result: A = {0,1,2,3,0,1,2,3,0,1,2,3}
```

# Synchronization

- Kernel-level synchronization
- Blocks must be independent
  - Can run in any order, concurrently or sequentially
  - Can't synchronize between blocks
- Some level of coordination can be achieved using atomic intrinsics → performance issues
- Threads in a block can synchronize using `__syncthreads` intrinsic

- Compilation tools are a part of CUDA SDK
- Compiler driver is called `nvcc`
- `nvcc` separates the code for host and device
  - Host code is compiled with regular C/C++ compiler
- `nvcc` or C/C++ can be used for linking
- Note that `nvcc` uses C++ front end to parse the program code
- For Linux environment NVIDIA provides a command-line debugger, `cuda-gdb`
  - For memory access violation checks there is also `cuda-memcheck`
- Profiling information can be gathered and visualized using `computeprof` tool

# Coalesced Memory Access

- Global memory access has very high latency
- Threads are executed in warps, memory operations are grouped in a similar fashion
  - Memory access is optimized for coalesced access where threads read from / write to successive memory locations
- Shared memory is better suited for more complicated data access

# Pinned Memory

- Normal malloc call returns a pointer to virtual memory → swapping, page faults
- Pinning is beneficial in some cases
  - Higher transfer speeds between host and device
  - Copying can be interleaved with kernel execution
- Pinning large memory areas can decrease overall system performance

# CUDA STREAMS

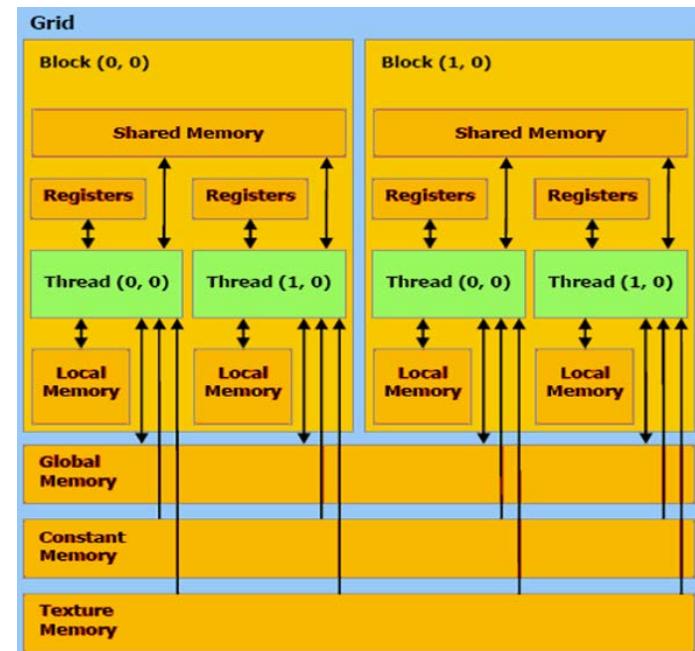
- Stream is a sequence of commands that execute in order
- Streams can be used to overlap memory copies and kernel execution
  - Copy-execute sequences are split into separate blocks that are associated to different streams

# ERROR CHECKING

- All runtime function calls return an error code
  - Important to check!
- Kernel launches do not return error codes
  - Have to use separate error checking routines
- Asynchronous operations can fail long after the call has returned → finding exact place of error can be difficult
  - For development phase extra synchronization may be useful

# SUMMARY

- CUDA programs consist of host and device code
- Device code is run parallel using threads organized into a grid of thread blocks
  - Kernel launch parameters
- Device and host have separate memories
  - Memory allocations and transfers



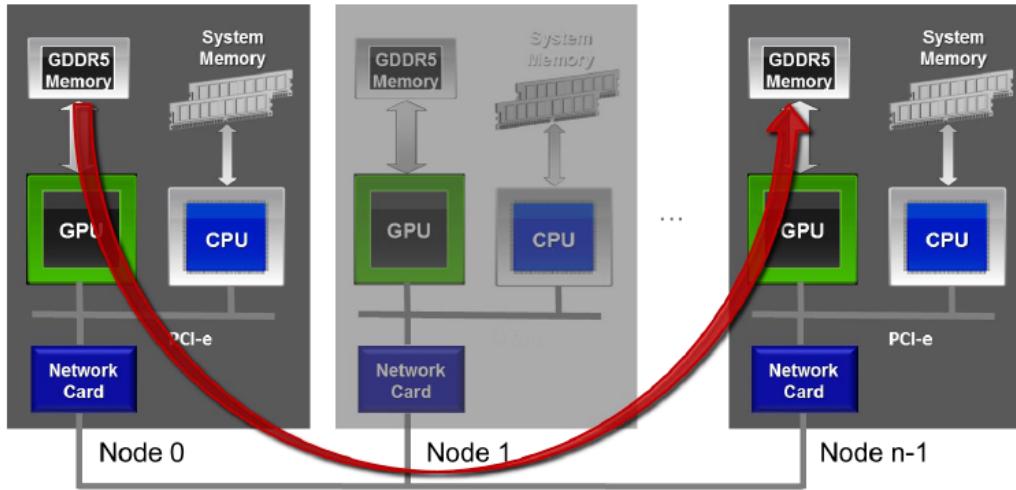
# CUDA + MPI + OpenMP

# OpenCL + MPI + OpenMP

Then we look at my notes that has runs and links to code run – on Discovery Cluster

<http://nuweb12.neu.edu/rc/wp-content/uploads/2014/11/CUDA-MPI.pdf>

# MPI+CUDA



```
//MPI rank 0  
MPI_Send(s_buf_d, size, MPI_CHAR, n-1, tag, MPI_COMM_WORLD);  
  
//MPI rank n-1  
MPI_Recv(r_buf_d, size, MPI_CHAR, 0, tag, MPI_COMM_WORLD, &stat);
```

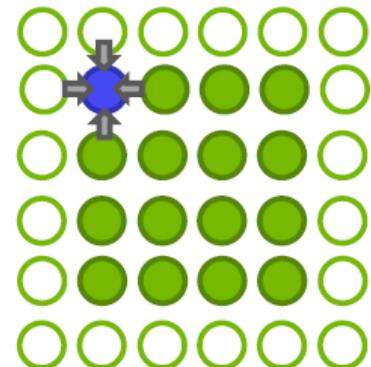
# EXAMPLE: JACOBI SOLVER - SINGLE GPU

While not converged

- Do Jacobi step:

```
for (int i=1; i < n-1; i++)  
    for (int j=1; j < m-1; j++)  
        u_new[i][j] = 0.0f - 0.25f*(u[i-1][j] + u[i+1][j]  
                                + u[i][j-1] + u[i][j+1])
```

- Swap \_new and
- Next iteration



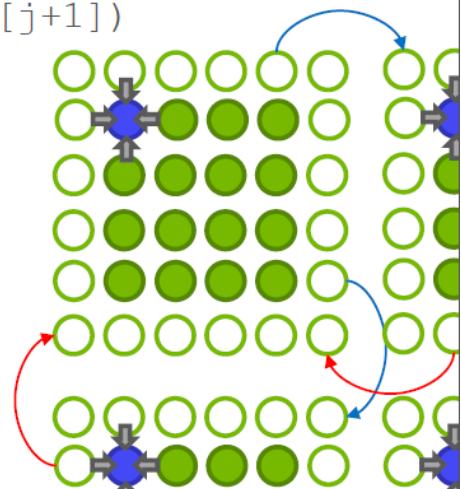
# EXAMPLE: JACOBI SOLVER - MULTI GPU

While not converged

- Do Jacobi step:

```
for (int i=1; i < n-1; i++)  
    for (int j=1; j < m-1; j++)  
        u_new[i][j] = 0.0f - 0.25f * (u[i-1][j] + u[i+1][j]  
   + u[i][j-1] + u[i][j+1])
```

- Exchange halo with 2 4 neighbor
- Swap \_new and
- Next iteration



# EXAMPLE: JACOBI SOLVER

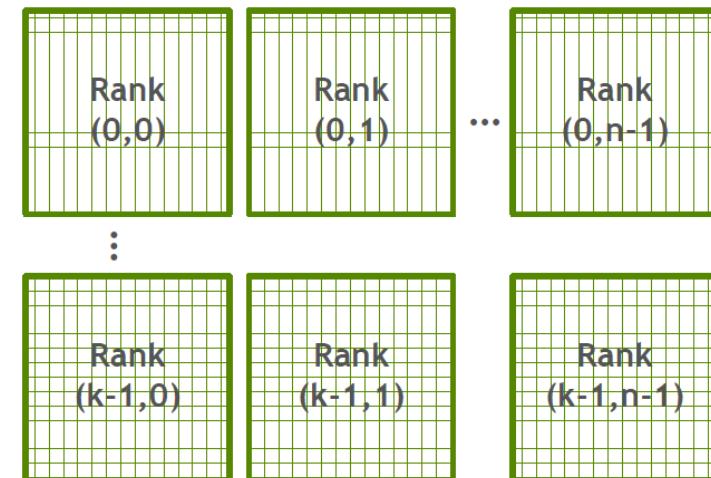
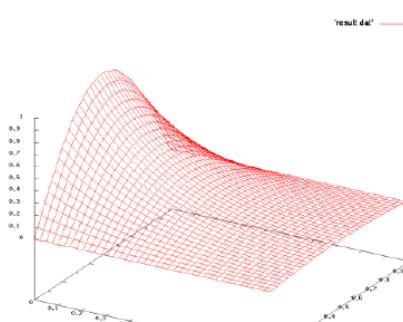
- Solves the 2D-Laplace equation on a rectangle

$$\Delta u(x, y) = 0 \quad \forall (x, y) \in \Omega \setminus \delta\Omega$$

- Dirichlet boundary conditions (constant values on boundaries)

$$u(x, y) = f(x, y) \in \delta\Omega$$

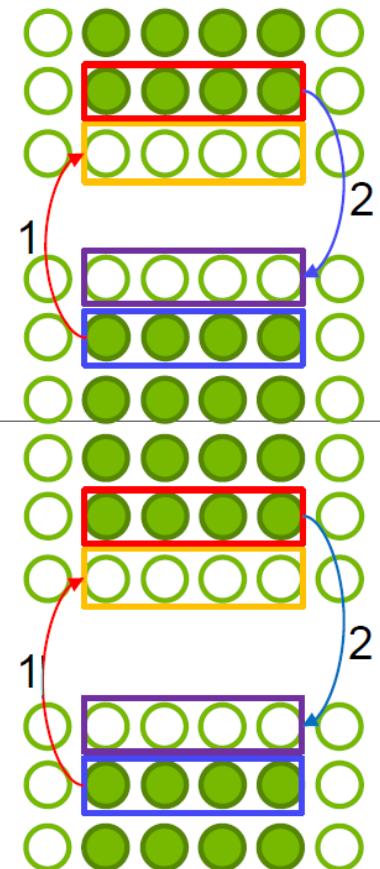
- 2D domain decomposition with  $n \times k$  domains



# EXAMPLE: JACOBI - TOP/BOTTOM HALO UPDATE

```
MPI_Sendrecv(u_new+offset_first_row, m-2, MPI_DOUBLE, t_nb, 0,  
             u_new+offset_bottom_boundary, m-2, MPI_DOUBLE, b_nb, 0,  
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```

```
MPI_Sendrecv(u_new+offset_last_row, m-2, MPI_DOUBLE, b_nb, 1,  
             u_new+offset_top_boundary, m-2, MPI_DOUBLE, t_nb, 1,  
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```



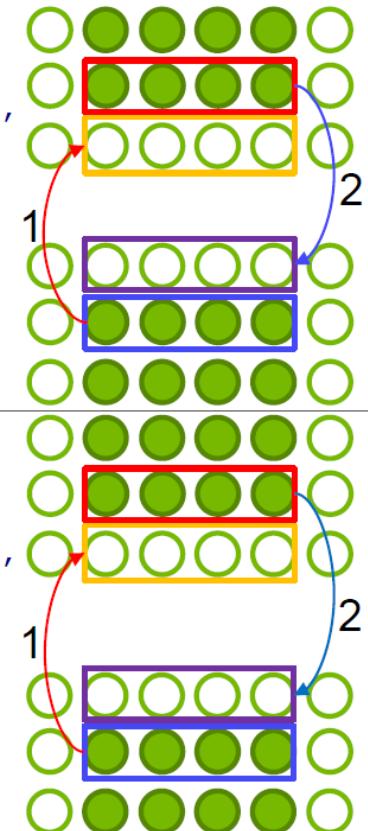
# EXAMPLE: JACOBI - TOP/BOTTOM HALO UPDATE

OpenACC

```
#pragma acc host_data use_device ( u_new ) {  
    MPI_Sendrecv( u_new+offset first row, m-2, MPI_DOUBLE, t_nb, 0,  
                  u_new+offset bottom boundary, m-2, MPI_DOUBLE, b_nb, 0,  
                  MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
  
    MPI_Sendrecv( u_new+offset last row, m-2, MPI_DOUBLE, b_nb, 1,  
                  u_new+offset top boundary, m-2, MPI_DOUBLE, t_nb, 1,  
                  MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
}
```

CUDA

```
MPI_Sendrecv( u_new_d+offset first row, m-2, MPI_DOUBLE, t_nb, 0,  
                  u_new_d+offset bottom boundary, m-2, MPI_DOUBLE, b_nb, 0,  
                  MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
  
MPI_Sendrecv( u_new_d+offset last row, m-2, MPI_DOUBLE, b_nb, 1,  
                  u_new_d+offset top boundary, m-2, MPI_DOUBLE, t_nb, 1,  
                  MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```



# EXAMPLE: JACOBI - LEFT/RIGHT HALO UPDATE

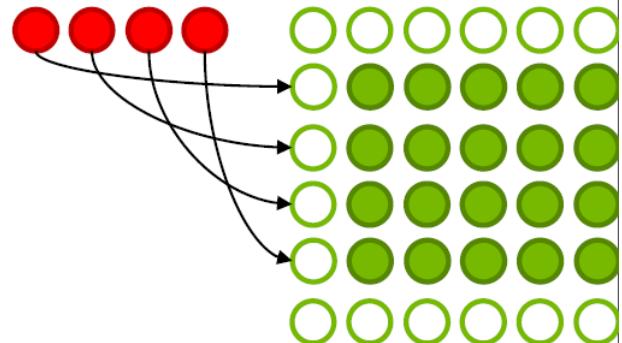
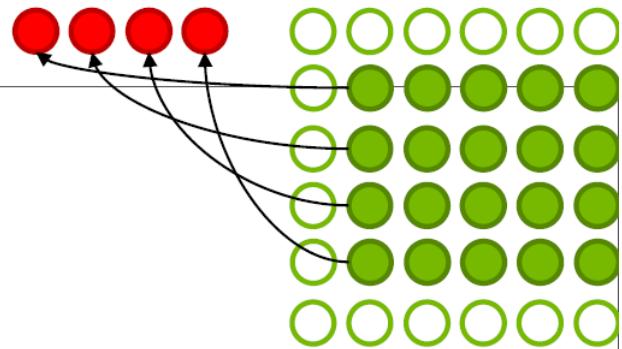
OpenACC

```
//right neighbor omitted

#pragma acc parallel loop present ( u_new, to_left )
for ( int i=0; i<n-2; ++i )
    to_left[i] = u_new[(i+1)*m+1];

#pragma acc host_data use_device ( from_left, to_left ) {
    MPI_Sendrecv( to_left, n-2, MPI_DOUBLE, l_nb, 0,
                  from_left, n-2, MPI_DOUBLE, l_nb, 0,
                  MPI_COMM_WORLD, MPI_STATUS_IGNORE );
}

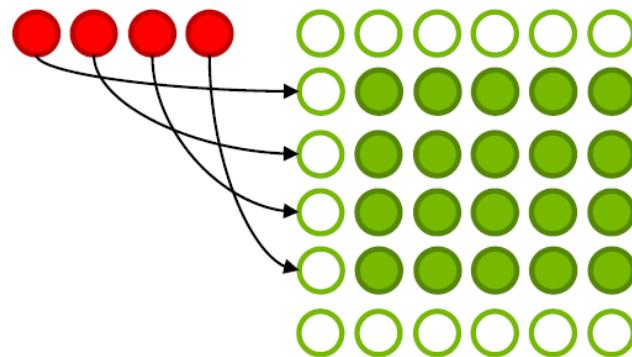
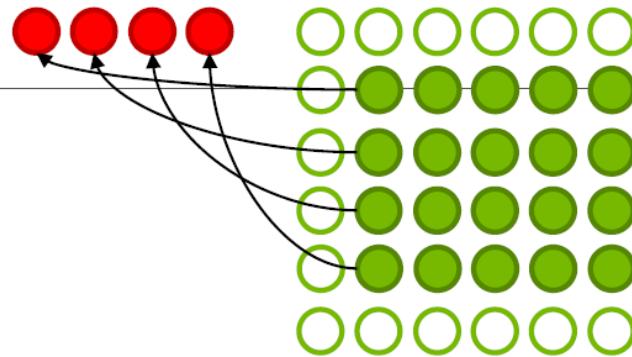
#pragma acc parallel loop present ( u_new, from_left )
for ( int i=0; i<n-2; ++i )
    u_new[(i+1)*m] = from_left[i];
```



# EXAMPLE: JACOBI - LEFT/RIGHT HALO UPDATE

CUDA

```
//right neighbor omitted  
pack<<<gs,bs,0,s>>>(to_left_d, u_new_d, n, m);  
cudaStreamSynchronize(s);  
  
MPI_Sendrecv( to_left_d, n-2, MPI_DOUBLE, l_nb, 0,  
              from_left_d, n-2, MPI_DOUBLE, l_nb, 0,  
              MPI_COMM_WORLD, MPI_STATUS_IGNORE );  
  
unpack<<<gs,bs,0,s>>>(u_new_d, from_left_d, n, m);
```



# EXAMPLE: JACOBI - TOP/BOTTOM HALO UPDATE - WITHOUT CUDA-AWARE MPI

OpenACC

```
#pragma acc update host( u_new[1:m-2], u_new[(n-2)*m+1:m-2] )
MPI_Sendrecv(u_new+offset_first_row, m-2, MPI_DOUBLE, t_nb, 0,
             u_new+offset_bottom_boundary, m-2, MPI_DOUBLE, b_nb, 0,
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
MPI_Sendrecv(u_new+offset_last_row, m-2, MPI_DOUBLE, b_nb, 1,
             u_new+offset_top_boundary, m-2, MPI_DOUBLE, t_nb, 1,
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);

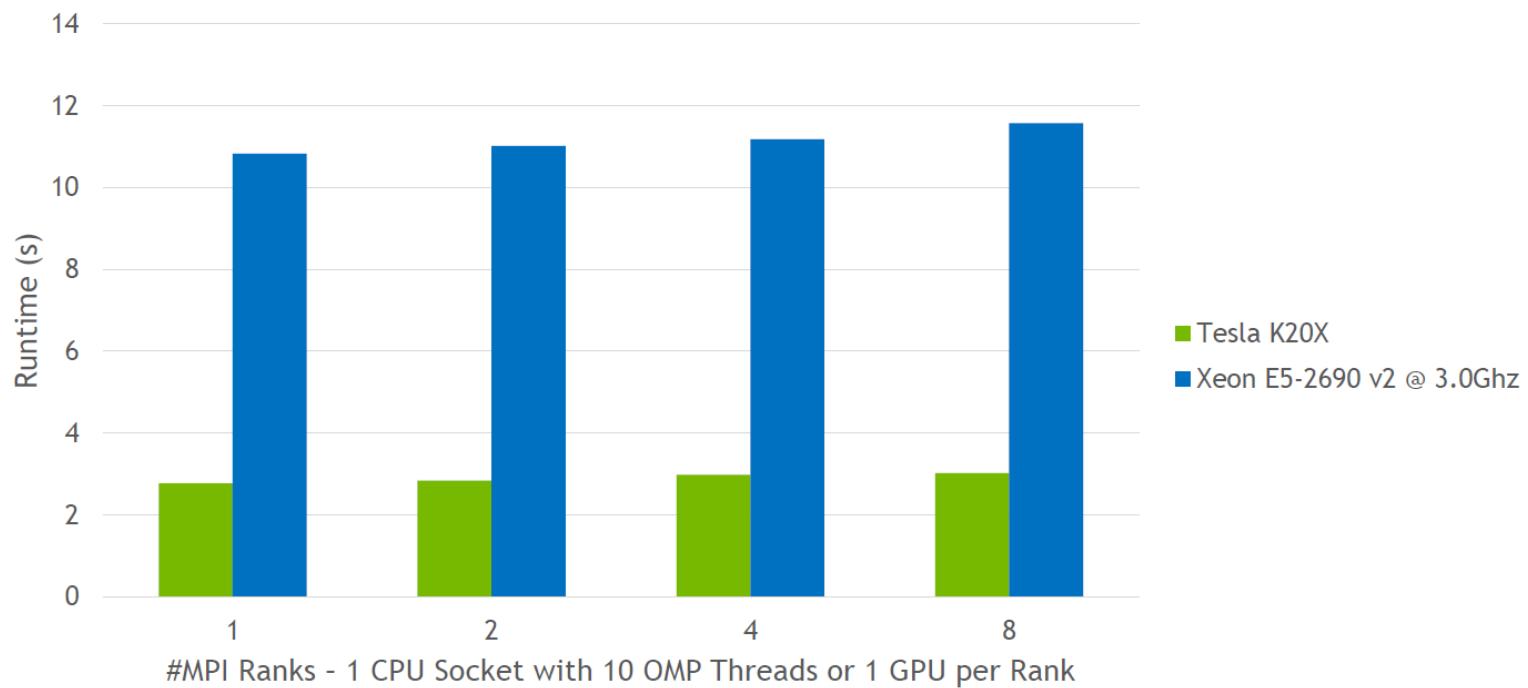
#pragma acc update device( u_new[0:m-2], u_new[(n-2)*m:m-2] )
//send to bottom and receive from top - top bottom omitted
```

CUDA

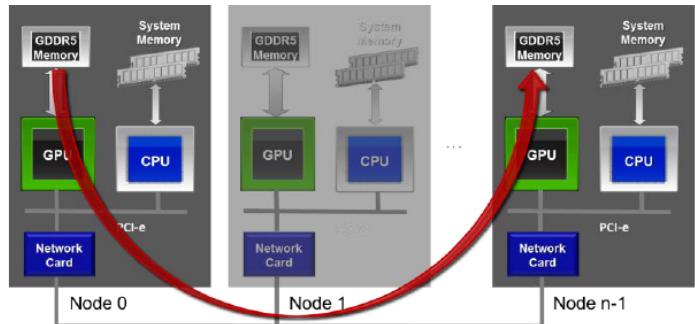
```
cudaMemcpy(u_new+1, u_new_d+1, (m-2)*sizeof(double), cudaMemcpyDeviceToHost);
MPI_Sendrecv(u_new+offset_first_row, m-2, MPI_DOUBLE, t_nb, 0,
             u_new+offset_bottom_boundary, m-2, MPI_DOUBLE, b_nb, 0,
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
cudaMemcpy(u_new_d, u_new, (m-2)*sizeof(double), cudaMemcpyDeviceToHost);
```

# JACOBI RESULTS (1000 STEPS)

## WEAK SCALING 4K X 4K PER PROCESS



# MPI+CUDA



## With UVA and CUDA-aware MPI No UVA and regular MPI

```
//MPI rank 0  
MPI_Send(s_buf_d, size, ...);
```

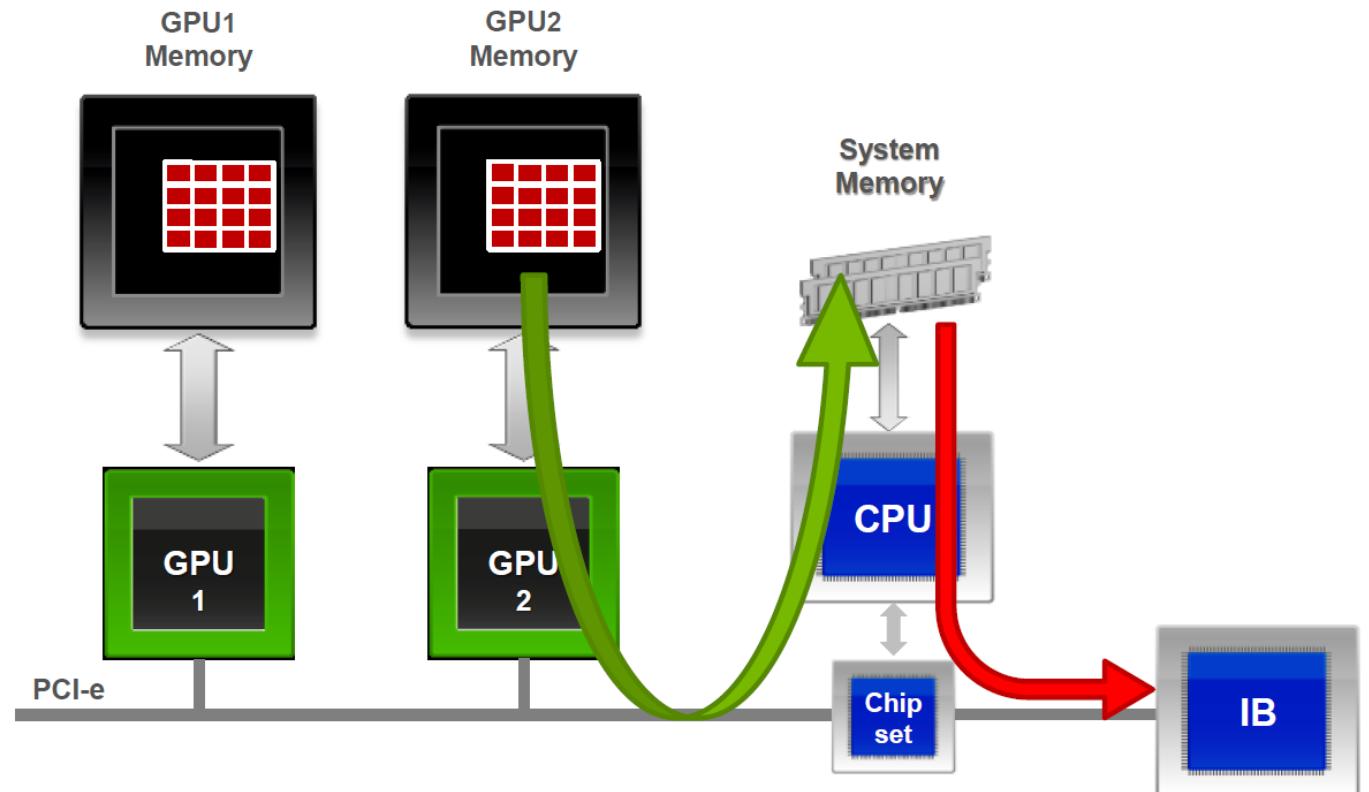
```
//MPI rank n-1  
MPI_Recv(r_buf_d, size, ...);
```

```
//MPI rank 0  
cudaMemcpy(s_buf_h, s_buf_d, size, ...);  
MPI_Send(s_buf_h, size, ...);
```

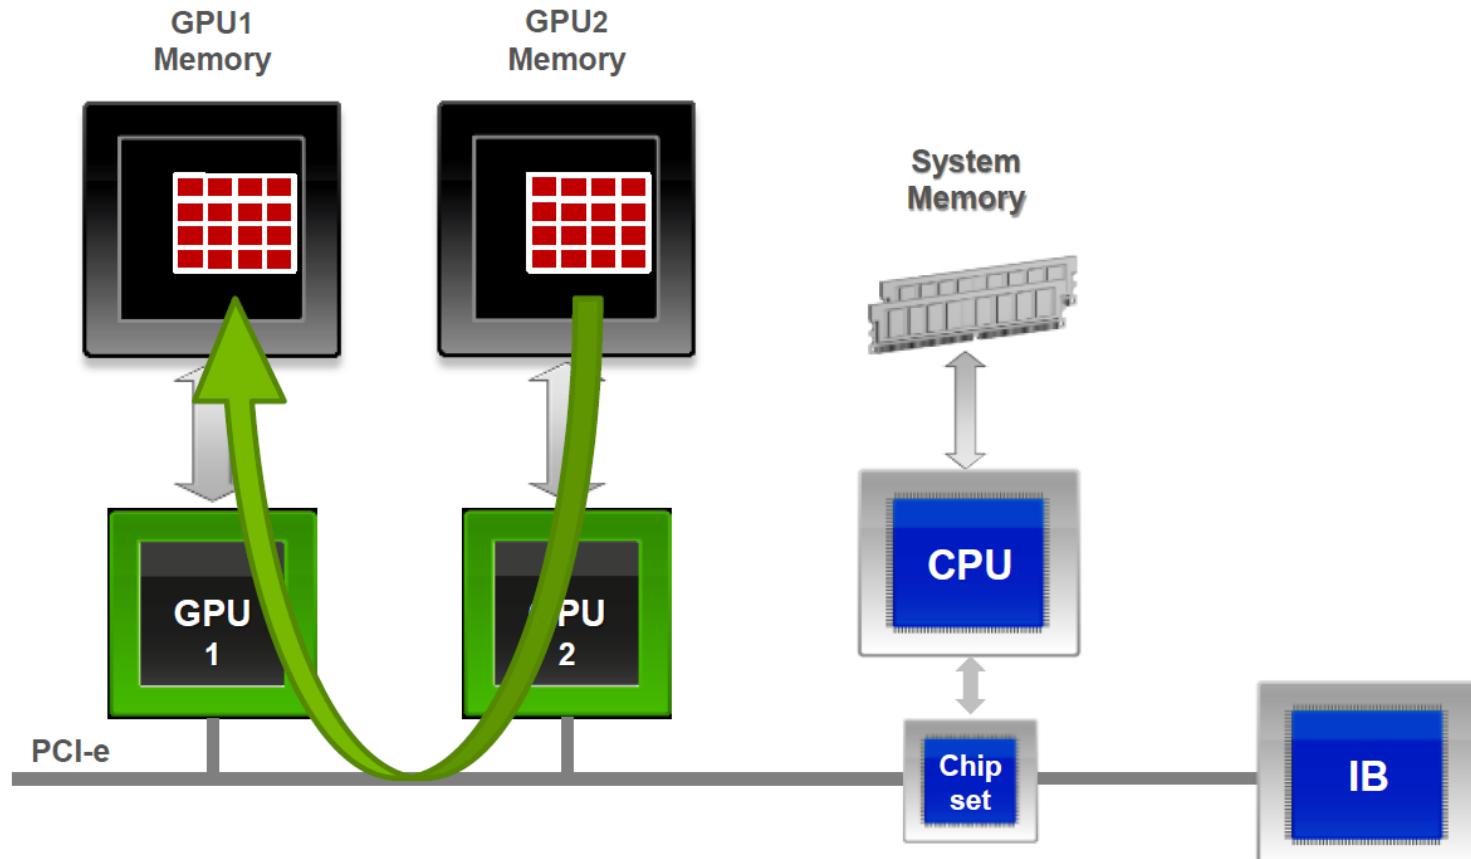
```
//MPI rank n-1  
MPI_Recv(r_buf_h, size, ...);  
cudaMemcpy(r_buf_d, r_buf_h, size, ...);
```

# NVIDIA GPUDIRECT™

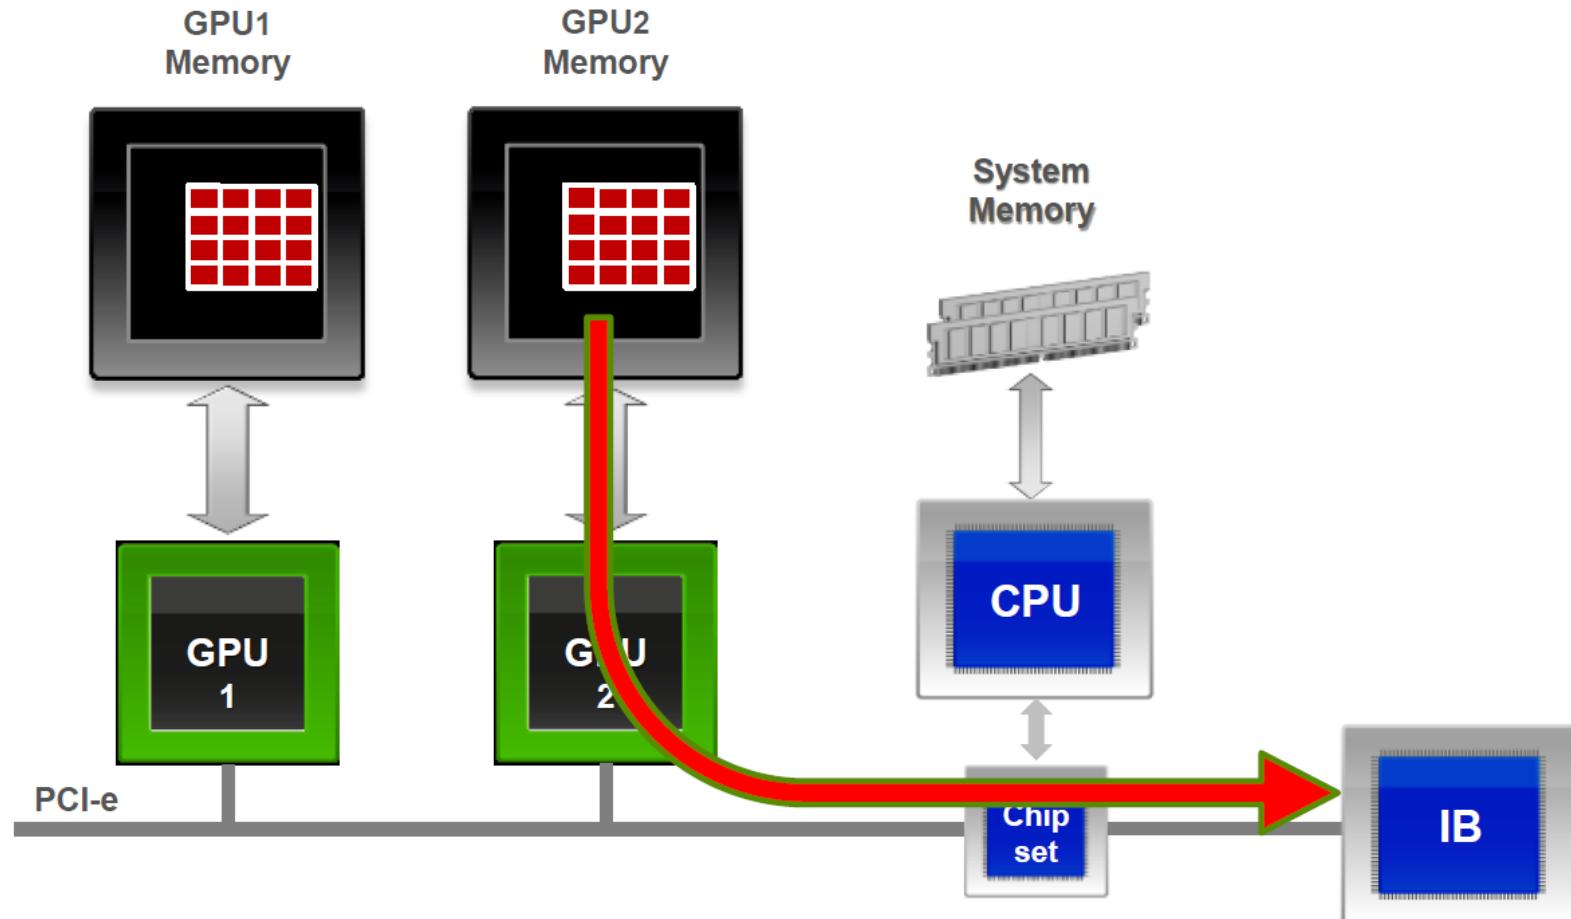
ACCELERATED COMMUNICATION WITH NETWORK & STORAGE DEVICES



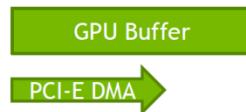
# NVIDIA GPUDIRECT™ PEER TO PEER TRANSFERS



# NVIDIA GPUDIRECT™ SUPPORT FOR RDMA



# CUDA-AWARE MPI



## MPI GPU TO REMOTE GPU GPUDIRECT SUPPORT FOR RDMA



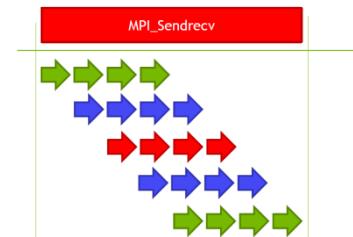
```
MPI_Send(s_buf_d,size,MPI_CHAR,1,tag,MPI_COMM_WORLD);
MPI_Recv(r_buf_d,size,MPI_CHAR,0,tag,MPI_COMM_WORLD,&stat);
```



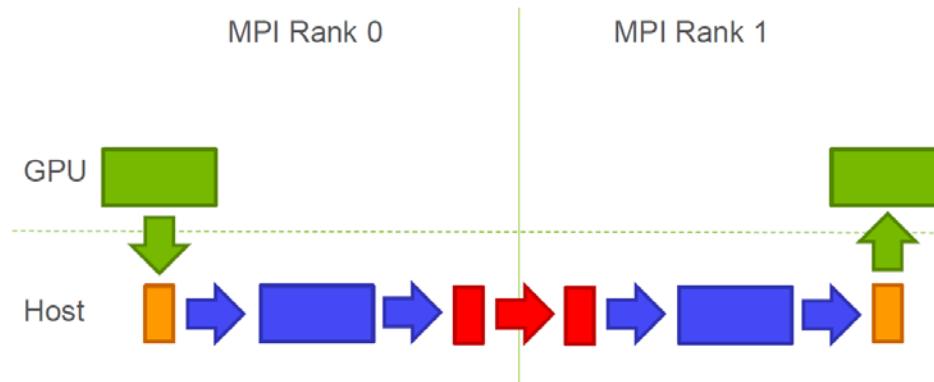
## MPI GPU TO REMOTE GPU WITHOUT GPUDIRECT



```
MPI_Send(s_buf_h,size,MPI_CHAR,1,tag,MPI_COMM_WORLD);
MPI_Recv(r_buf_h,size,MPI_CHAR,0,tag,MPI_COMM_WORLD,&stat);
```



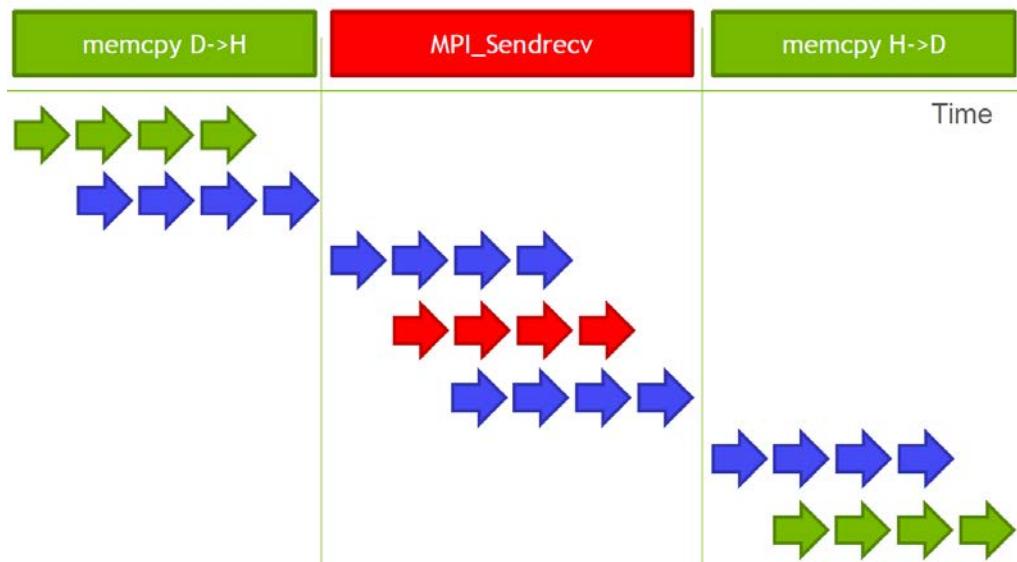
## REGULAR MPI GPU TO REMOTE GPU



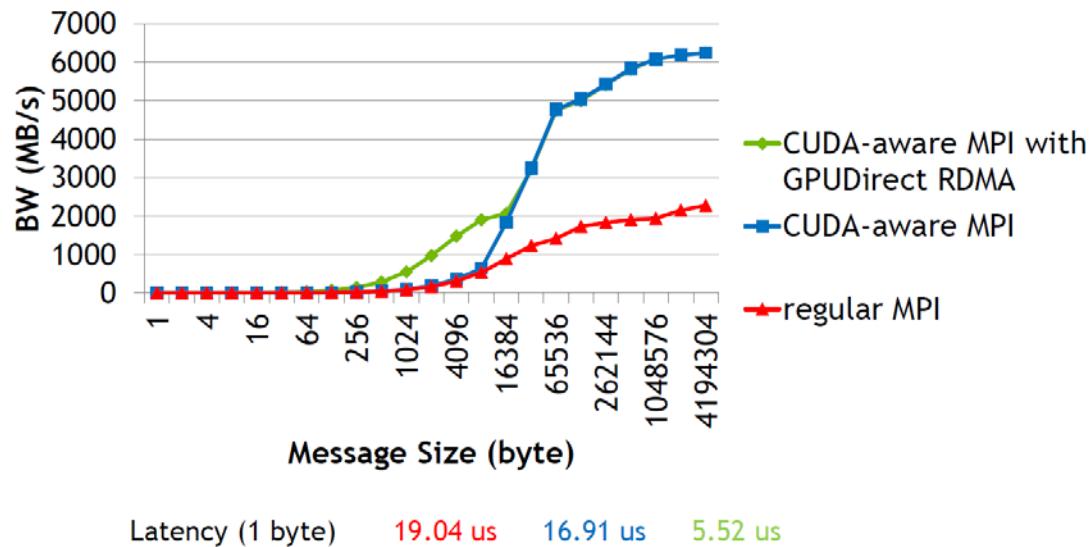
```
cudaMemcpy(s_buf_h,s_buf_d,size,cudaMemcpyDeviceToHost);
MPI_Send(s_buf_h,size,MPI_CHAR,1,tag,MPI_COMM_WORLD);

MPI_Recv(r_buf_h,size,MPI_CHAR,0,tag,MPI_COMM_WORLD,&stat);
cudaMemcpy(r_buf_d,r_buf_h,size,cudaMemcpyHostToDevice);
```

## REGULAR MPI GPU TO REMOTE GPU



## PERFORMANCE RESULTS TWO NODES - EXAMPLE



# GPU ACCELERATION OF LEGACY MPI APPLICATION

Typical legacy application

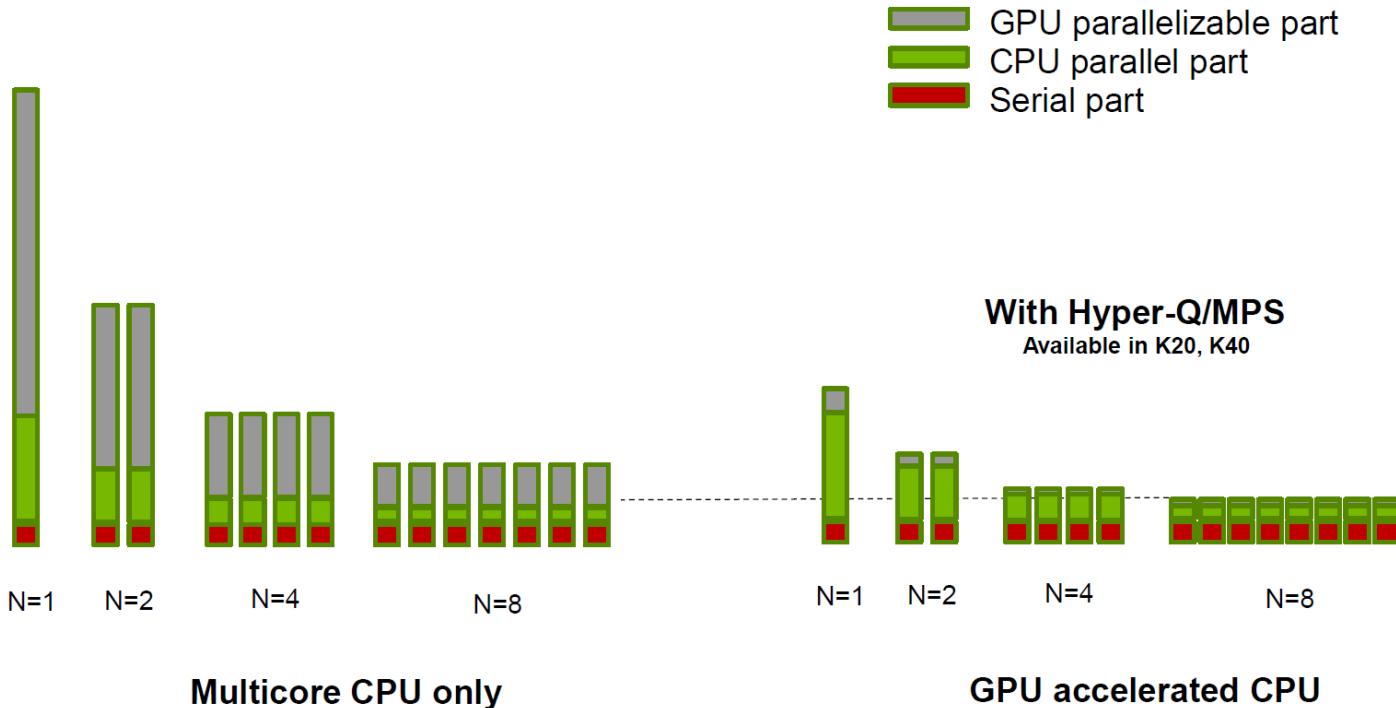
- MPI parallel
- Single or few threads per MPI rank (e.g. OpenMP)

Running with multiple MPI ranks per node

GPU acceleration in phases

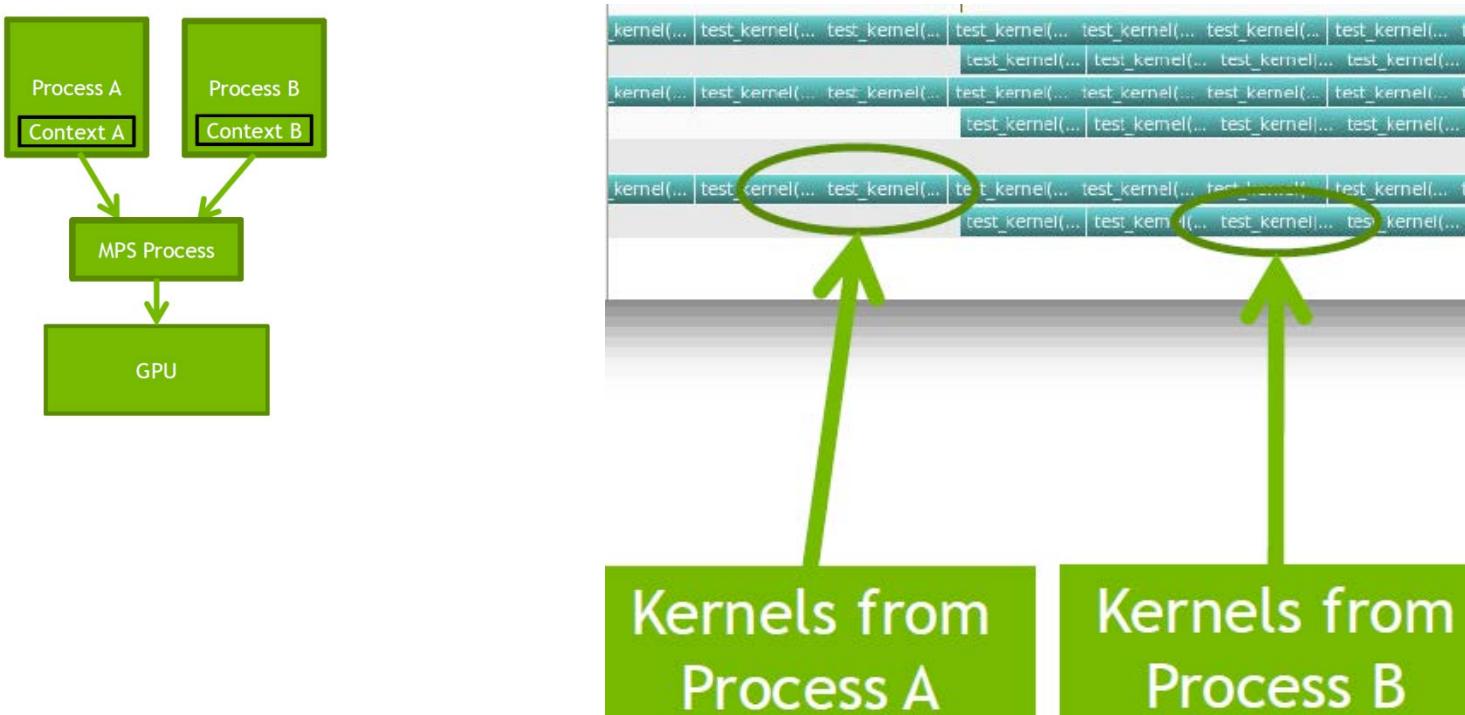
- Proof of concept prototype, ..
- Great speedup at kernel level

Application performance misses expectations



## PROCESSES SHARING GPU WITH MPS / HYPER-Q:

- MAXIMUM OVERLAP
- Enables overlap between copy and compute of different processes
- Sharing the GPU between multi-MPI ranks increases GPU utilization

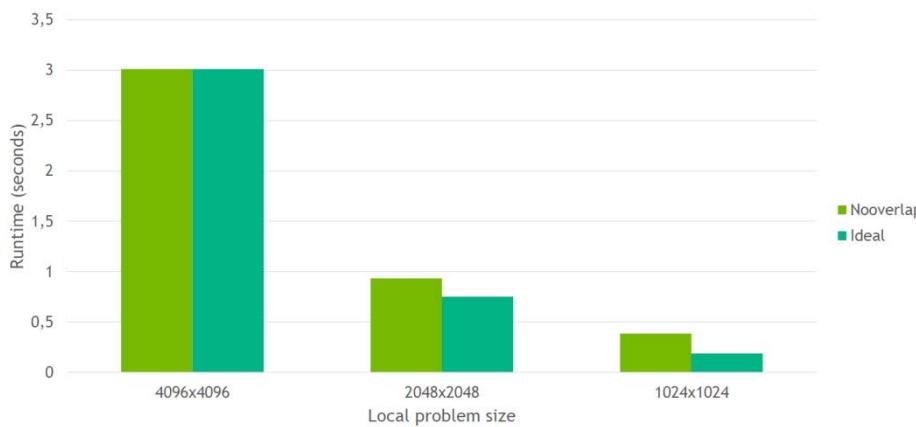


# BEST PRACTICE: USE NONE-BLOCKING MPI

BLOCKING  
NONE-BLOCKING

```
#pragma acc host_data use_device ( u_new ) {  
  
    MPI_Sendrecv(u_new+offset_first_row, m-2, MPI_DOUBLE, t_nb, 0,  
                u_new+offset_bottom_boundary, m-2, MPI_DOUBLE, b_nb, 0,  
                MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
  
    MPI_Sendrecv(u_new+offset_last_row, m-2, MPI_DOUBLE, b_nb, 1,  
                u_new+offset_top_boundary, m-2, MPI_DOUBLE, t_nb, 1,  
                MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
  
}  
  
MPI_Request t_b_req[4];  
  
#pragma acc host_data use_device ( u_new ) {  
  
    MPI_Irecv(u_new+offset_top_boundary,m-2,MPI_DOUBLE,t_nb,MPI_STATUS_IGNORE,t_b_req+0);  
    MPI_Irecv(u_new+offset_bottom_boundary,m-2,MPI_DOUBLE,b_nb,MPI_STATUS_IGNORE,t_b_req+1);  
    MPI_Isend(u_new+offset_last_row,m-2,MPI_DOUBLE,b_nb,MPI_STATUS_IGNORE,t_b_req+2);  
    MPI_Isend(u_new+offset_first_row,m-2,MPI_DOUBLE,t_nb,1,MPI_COMM_WORLD,t_b_req+3);  
  
}  
  
MPI_Waitall(4, t_b_req, MPI_STATUSES_IGNORE);
```

Gives MPI more  
opportunities to build  
efficient pipelines

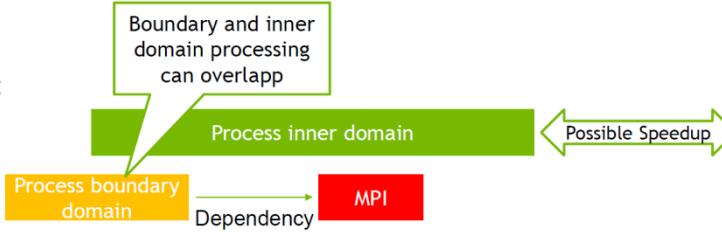


# OVERLAPPING COMMUNICATION AND COMPUTATION

No Overlap



Overlap



CUDA

```
process_boundary_and_pack<<<gs_b,bs_b,0,s1>>>(u_new_d,u_d,to_left_d,to_right_d,n,m);
```

```
process_inner_domain<<<gs_id,bs_id,0,s2>>>(u_new_d, u_d,to_left_d,to_right_d,n,m);
```

```
cudaStreamSynchronize(s1); //wait for boundary
```

```
MPI_Request req[8];
```

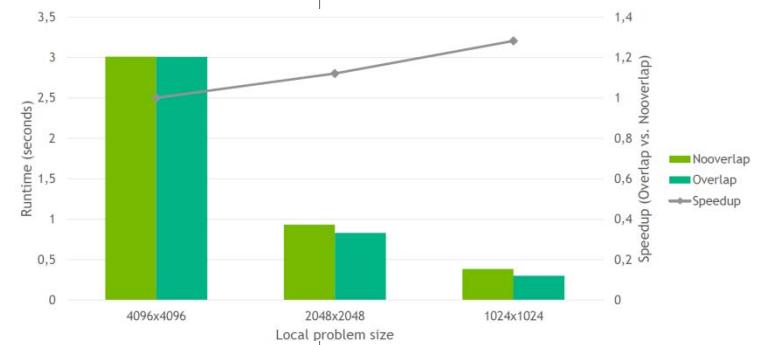
```
//Exchange halo with left, right, top and bottom neighbor
```

```
MPI_Waitall(8, req, MPI_STATUSES_IGNORE);
```

```
unpack<<<gs_s,bs_s>>>(u_new_d, from_left_d, from_right_d, n, m);
```

```
cudaDeviceSynchronize(); //wait for iteration to finish
```

```
#pragma acc parallel loop present ( u_new, u, to_left, to_right ) async(1)
for ( ... )
    //Process boundary and pack to_left and to_right
#pragma acc parallel loop present ( u_new, u ) async(2)
for ( ... )
    //Process inner domain
#pragma acc wait(1)                                //wait for boundary
MPI_Request req[8];
#pragma acc host_data use_device ( from_left, to_left, form_right, to_right, u_new ) {
    //Exchange halo with left, right, top and bottom neighbor
}
MPI_Waitall(8, req, MPI_STATUSES_IGNORE);
#pragma acc parallel loop present ( u_new, from_left, from_right )
for ( ... )
    //unpack from_left and from_right
#pragma acc wait                                //wait for iteration to finish
```



# MPI AND UNIFIED MEMORY

- Unified Memory support for CUDA-aware MPI needs changes to the MPI implementations
  - Check with your MPI implementation of choice for their plans
  - It might work in some situations but it is not supported
- Unified Memory and regular MPI
  - Require unmanaged staging buffers
    - Regular MPI has no knowledge of managed memory
    - CUDA 6 managed memory does not play well with RDMA protocols

# HANDLING MULTI GPU NODES

- Multi GPU nodes and GPU-affinity:

- Use local rank:

```
int local_rank = //determine local rank  
int num_devices = 0;  
cudaGetDeviceCount(&num_devices);  
cudaSetDevice(local_rank % num_devices);
```

- Use exclusive process mode + cudaSetDevice(0)

- How to determine local rank:

- Rely on process placement (with one rank per GPU)

```
int rank = 0;  
MPI_Comm_rank(MPI_COMM_WORLD,&rank);  
int num_devices = 0;  
cudaGetDeviceCount(&num_devices); // num_devices == ranks per node  
int local_rank = rank % num_devices;
```

- Use environment variables provided by MPI launcher

- e.g for OpenMPI

```
int local_rank = atoi(getenv("OMPI_COMM_WORLD_LOCAL_RANK"));
```

## TOOLS FOR MPI+CUDA APPLICATIONS

- Memory Checking `cuda-memcheck`
- Debugging `cuda-gdb`
- Profiling `nvprof` and NVIDIA Visual Profiler

## MEMORY CHECKING WITH CUDA-MEMCHECK

- Cuda-memcheck is a functional correctness checking suite similar to the valgrind memcheck tool
- Can be used in a MPI environment

```
mpiexec -np 2 cuda-memcheck ./myapp <args>
```

- Problem: output of different processes is interleaved

- Use save, log-file command line options and launcher script

```
#!/bin/bash  
  
LOG=$1.$OMPI_COMM_WORLD_RANK  
#LOG=$1.$MV2_COMM_WORLD_RANK  
cuda-memcheck --log-file $LOG.log --save $LOG.memcheck $*  
  
mpiexec -np 2 cuda-memcheck-script.sh ./myapp <args>
```

## DEBUGGING MPI+CUDA APPLICATIONS

### USING CUDA-GDB WITH MPI APPLICATIONS

- You can use cuda-gdb just like gdb with the same tricks
  - For smaller applications, just launch xterms and cuda-gdb
- ```
> mpiexec -x -np 2 xterm -e cuda-gdb ./myapp <args>
```

## DEBUGGING MPI+CUDA APPLICATIONS

### CUDA-GDB ATTACH

- CUDA 5.0 and forward have the ability to attach to a running process

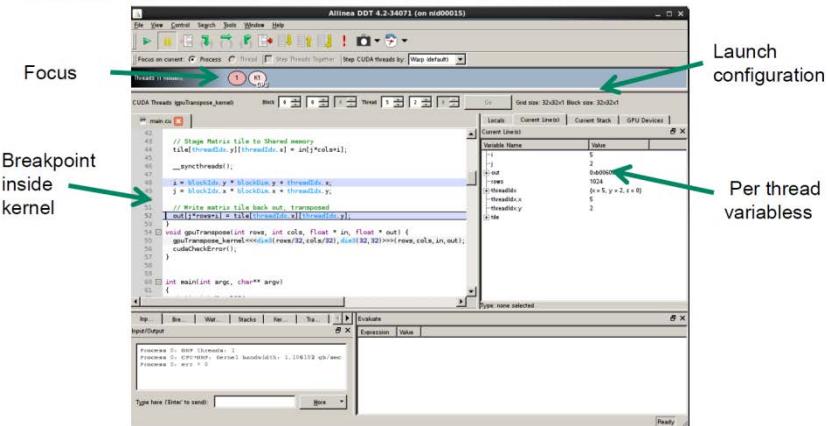
```
if ( rank == 0 ) {  
    int i=0;  
    printf("rank %d: pid %d on %s ready for attach\n", rank, getpid(), name);  
    while (0 == i) {  
        sleep(5);  
    }  
}  
  
> mpiexec -np 2 ./jacobi_mpi+cuda  
rank 0: pid 30034 on ge107 ready for attach  
> ssh ge107  
cuda-gdb --pid 30034
```

## DEBUGGING MPI+CUDA APPLICATIONS

### THIRD PARTY TOOLS

- Allinea DDT debugger
- Totalview

## DDT: THREAD LEVEL DEBUGGING



# PROFILING MPI+CUDA APPLICATIONS

## USING NVPROF+NVVP

### 3 Usage modes:

- Embed pid in output filename

```
mpirun -np 2 nvprof --output-profile profile.out.%p
```

- Only save the textual output

```
mpirun -np 2 nvprof --log-file profile.out.%p
```

- Collect profile data on all processes that run on a node

```
nvprof --profile-all-processes -o profile.out.%p
```

# PROFILING MPI+CUDA APPLICATIONS

## THIRD PARTY TOOLS

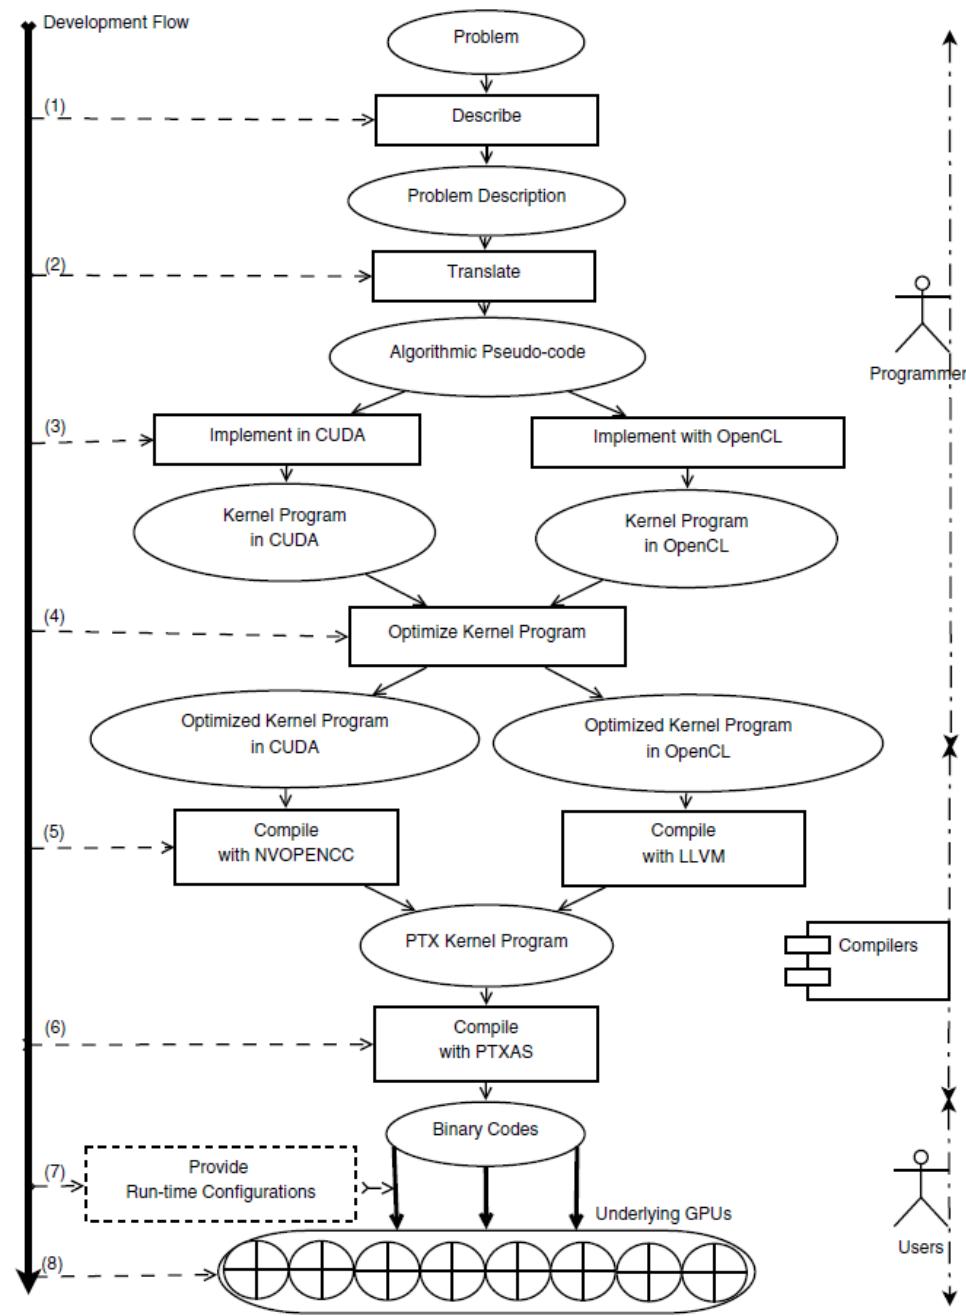
- Multiple parallel profiling tools are CUDA aware
  - Score-P
  - Vampir
  - Tau
- These tools are good for discovering MPI issues as well as basic CUDA performance inhibitors

# OVERLAPPING COMMUNICATION AND COMPUTATION - TIPS AND TRICKS

- CUDA-aware MPI might use the default stream
  - Allocate stream with the non-blocking flag (cudaStreamNonBlocking)
- In case of multiple kernels for boundary handling the kernel processing the inner domain might sneak in
  - Use single stream or events for inter stream dependencies via cudaStreamWaitEvent (#pragma acc wait async) - disables overlapping of boundary and inner domain kernels
  - Use high priority streams for boundary handling kernels - allows overlapping of boundary and inner domain kernels
- As of CUDA 6.0 GPUDirect P2P in multi process can overlap disable it for older releases

## CONCLUSIONS

- Using MPI as abstraction layer for Multi GPU programming allows multi GPU programs to scale beyond a single node
  - CUDA-aware MPI delivers ease of use, reduced network latency and increased bandwidth
- All NVIDIA tools are usable and third party tools are available
- Multiple CUDA-aware MPI implementations available
  - OpenMPI, MVAPICH2, Cray, IBM Platform MPI



# Debugging MPI with MPE

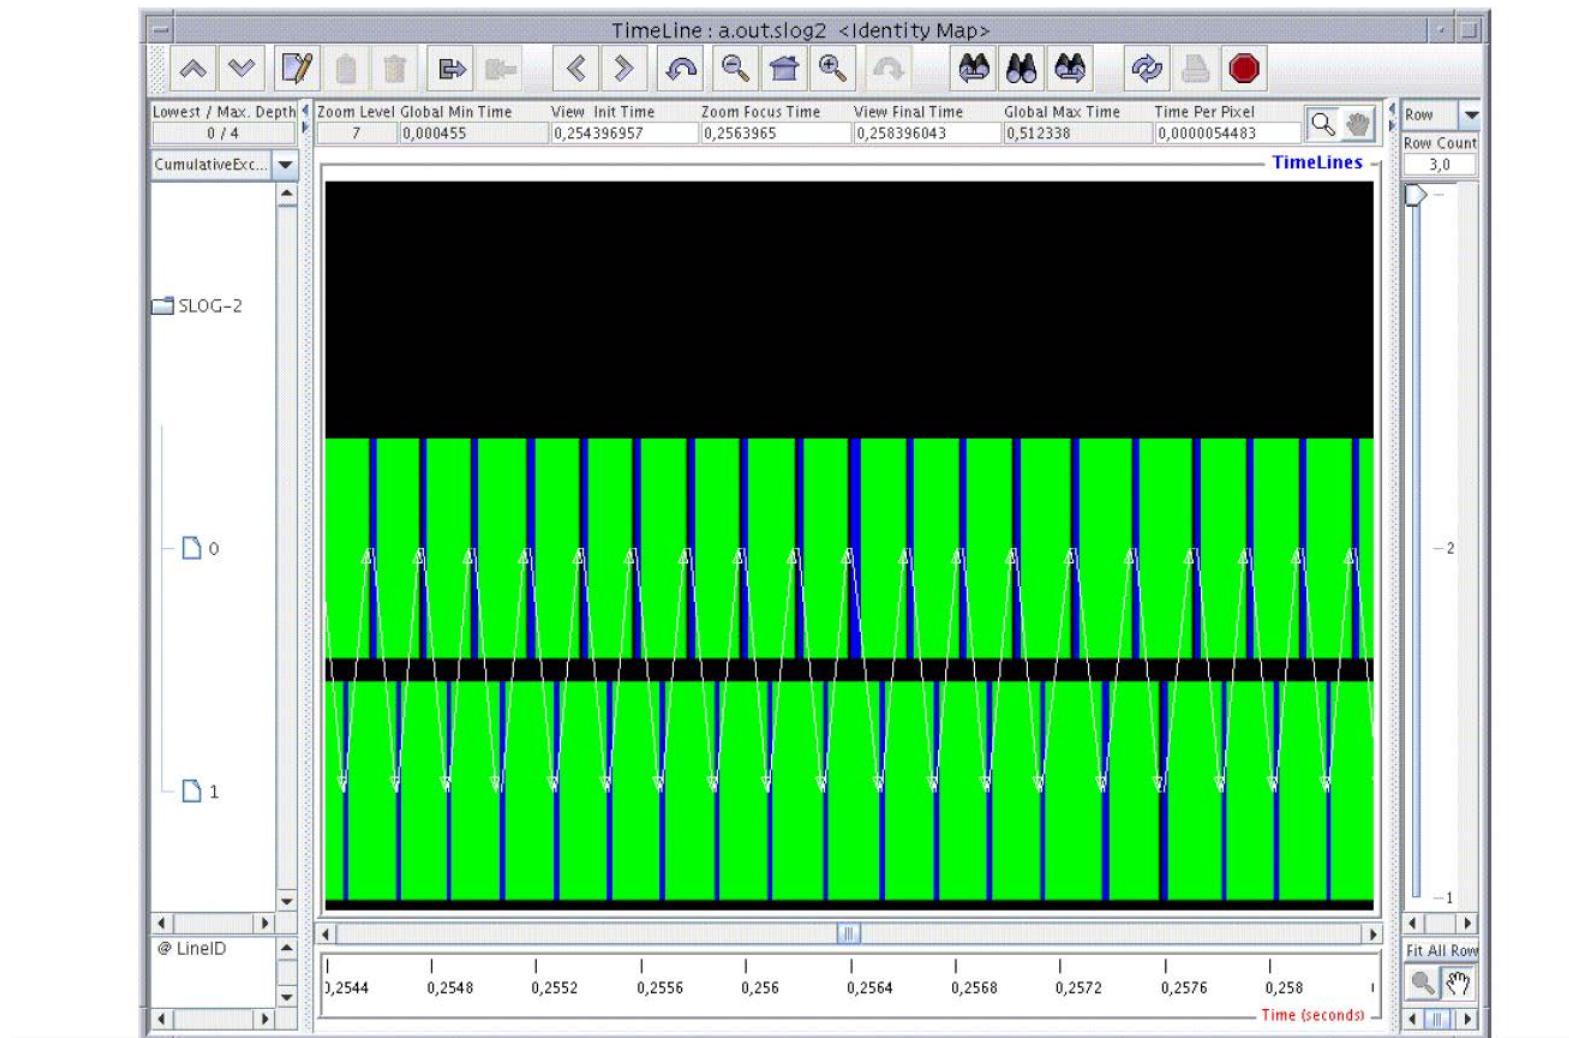
## MPI Parallel Environment (MPE)

- Software package for MPI programmers
- Provides users with a number of useful tool
  - visualisation, log converters, tracers
- Documentation
  - <http://www-unix.mcs.anl.gov/perfvis/download/index.htm>
- Compile the MPI program with the mpecc wrapper
  - **-mpilog**: Automatic MPI and MPE user-defined states logging
  - **-mpitrace**: Trace MPI program with printf
  - **-mpianim**: Animate MPI program in real-time.
  - ...
- Log file formats
  - ALOG (ASCII), CLOG (BINARY) maintained for compatibility reasons
  - SLOG = Scalable log

# MPI Program Tracing

- `mpecc -mpilog mpi_latency.c`
- `mpirun –np 2 a.out`
  - Produces a `a.out.clog` trace file
  - Convert to SLOG format using `clogToslog2` program
- `mpirun -np 2 ./a.out`
  - `-MPDENV -MPE_LOG_FORMAT=SLOG`
  - Produces a `a.out.slog2` file
- Open and visualize the **slogs** file with Jumpshot

# Jumpshot Latency Program Snapshot



Let's look at the parallel matrix multiplication example with MPE and Jumpshot on Discovery Cluster

- Performance Counters
- Profiling
  - PAPI
  - TAU
  - HPCToolkit
- Performance Application Programming  
Interface
- To design, standardize and implement a portable and efficient API
- C and Fortran Interface

```
[nroy@discovery2 utils]$ ./papi_avail  
Available events and hardware information.
```

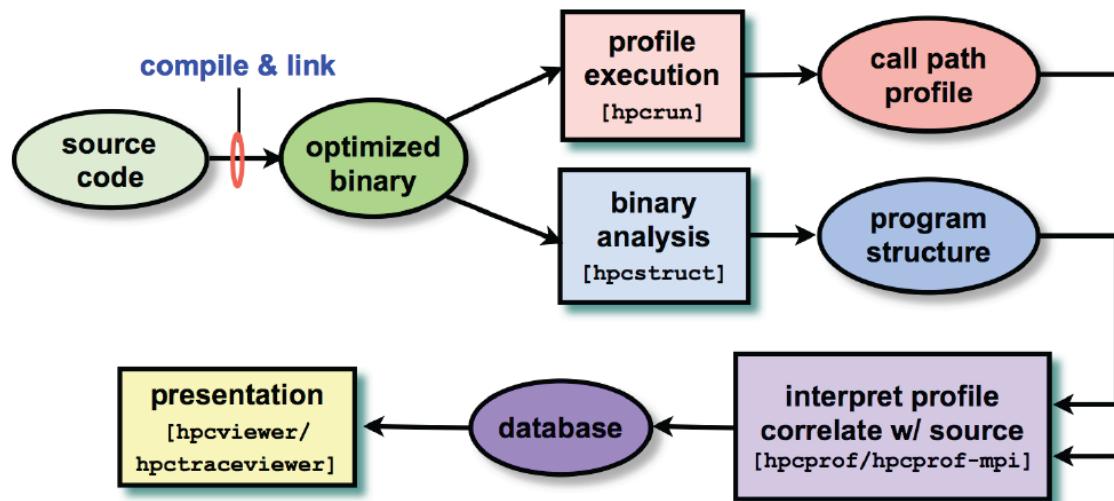
```
-----  
PAPI Version      : 5.3.0.0  
Vendor string and code : GenuineIntel (1)  
Model string and code : Intel(R) Xeon(R) CPU E5-2670 0 @ 2.60GHz (45)  
CPU Revision      : 7.000000  
CPUID Info        : Family: 6 Model: 45 Stepping: 7  
CPU Max Megahertz : 2599  
CPU Min Megahertz : 2599  
Hdw Threads per core : 2  
Cores per Socket   : 8  
Sockets            : 2  
NUMA Nodes         : 2  
CPUs per Node      : 16  
Total CPUs         : 32  
Running in a VM    : no  
Number Hardware Counters : 11  
Max Multiplex Counters : 32  
-----
```

Name	Code	Avail	Deriv	Description (Note)
PAPI_L1_DCM	0x80000000	Yes	No	Level 1 data cache misses
PAPI_L1_ICM	0x80000001	Yes	No	Level 1 instruction cache misses
PAPI_L2_DCM	0x80000002	Yes	Yes	Level 2 data cache misses
PAPI_L2_ICM	0x80000003	Yes	No	Level 2 instruction cache misses
PAPI_L3_DCM	0x80000004	No	No	Level 3 data cache misses
PAPI_L3_ICM	0x80000005	No	No	Level 3 instruction cache misses
PAPI_L1_TCM	0x80000006	Yes	Yes	Level 1 cache misses
PAPI_L2_TCM	0x80000007	Yes	No	Level 2 cache misses
PAPI_L3_TCM	0x80000008	Yes	No	Level 3 cache misses
.....				
.....				
.....				
PAPI_BTAC_M	0x8000001b	No	No	Branch target address cache misses
instructions				
PAPI_FNV_INS	0x80000065	No	No	Floating point inverse instructions
PAPI_FP_OPS	0x80000066	Yes	Yes	Floating point operations
PAPI_SP_OPS	0x80000067	Yes	Yes	Floating point operations; optimized to count scaled single precision vector operations
PAPI_DP_OPS	0x80000068	Yes	Yes	Floating point operations; optimized to count scaled double precision vector operations
PAPI_VEC_SP	0x80000069	Yes	Yes	Single precision vector/SIMD instructions
PAPI_VEC_DP	0x8000006a	Yes	Yes	Double precision vector/SIMD instructions
PAPI_REF_CYC	0x8000006b	Yes	No	Reference clock cycles

```
-----  
Of 108 possible events, 50 are available, of which 17 are derived.
```

```
avail.c          PASSED  
[nroy@discovery2 utils]$ pwd  
/shared/apps/hpctoolkit/papi-5.3.0/src/utils  
[nroy@discovery2 utils]$
```

- <http://hpctoolkit.org>
- Measurement and analysis of program performance
- Using statistical sampling of timers and hardware performance counters
- Platforms supported: Linux X86\_64, Linux-x86, Linux-Power, Cray XT/XE/XK, IBM Blue Gene/Q, Blue Gene/P



## hpcrun

collecting calling-context-sensitive performance measurements

## hpcstruct

Analyzing application binaries and information between binaries and source codes

## hpcprof

correlating the result with source code

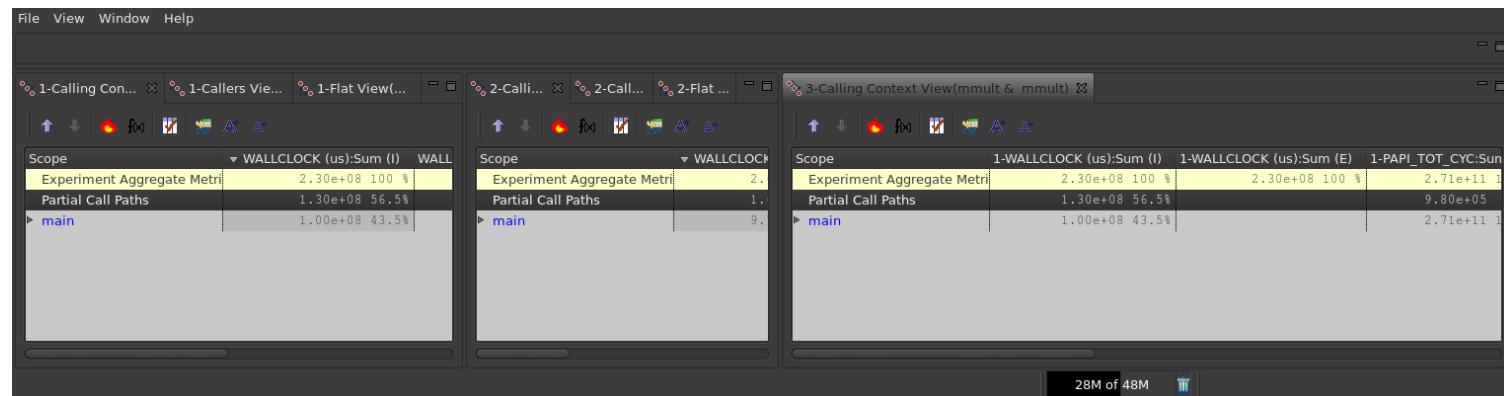
## hpcviewer

graphical user interface that presents a hierarchical, time-centric view of a program execution

Weak/strong scaling

Example: HPCToolkit use in determining scalability of programs (and assigning blame on a function to function basis). Since now we are only interested in execution time and number of cycles taken, we use the following counters.

Counter name	Description	Period
PAPI_TOT_CYC	Total cycles	10000
WALLCLOCK	Wall clock time used by the process in microseconds	100000



- Let go to Discovery Cluster and run a example.

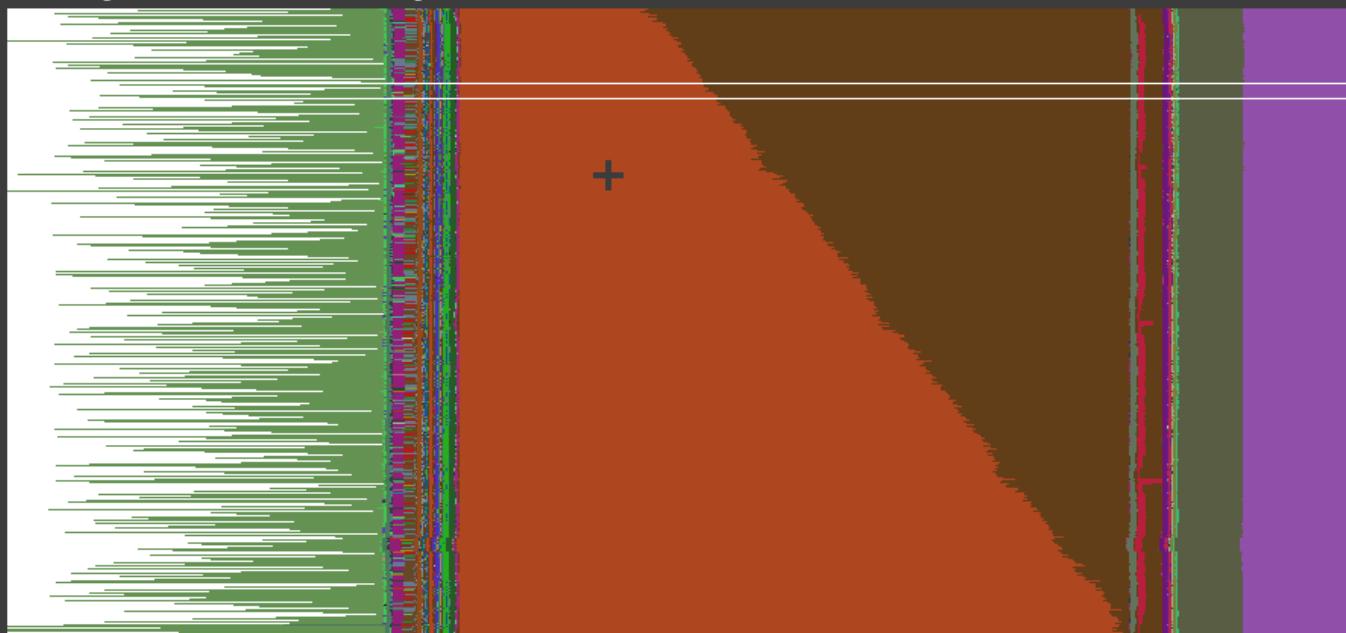
```
mpicc -g -O3 -o cpi_debug cpi.c -gdwarf-2  
hpcstruct ./cpi_debug  
bsub < bsubmit_debug.bash  
hpcprof-mpi -S cpi_debug.hpcstruct -l ./*' hpctoolkit-cpi_debug-measurements  
hpctraceviewer hpctoolkit-cpi_debug-database  
hpcviewer hpctoolkit-cpi_debug-database
```

File View Window Help

Trace View



Time Range: [0.829s ,18.335s] Rank Range: [0.0,1023.0] Cross Hair: (8.588s, 271.0)

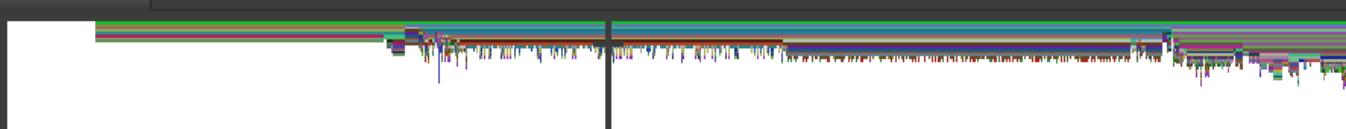


Call Path

8

- main
- amrGodunov()
- AMR::setupForFixedHierarchyRun()
- inlined from AMR.cpp: 259
- AMRLevelPolytropicGas::initialGrid()
- AMRLevelPolytropicGas::levelSetup()
- LevelGodunov::define(DisjointBox)
- inlined from LevelGodunov.cpp: 5
- PiecewiseLinearFillPatch::define(D)
- LayoutIterator::getToRealBox()

Depth View



Mini Map



Summary View



27M of 206M

THANK YOU  
QUESTIONS?